

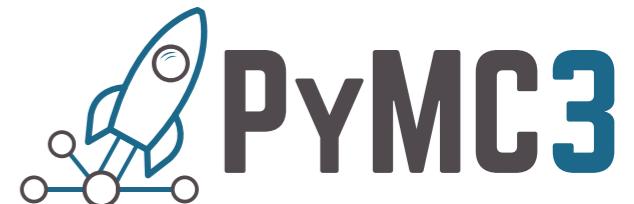
PRECISION WORKSHOP

Practical Bayesian Computation with PyMC3

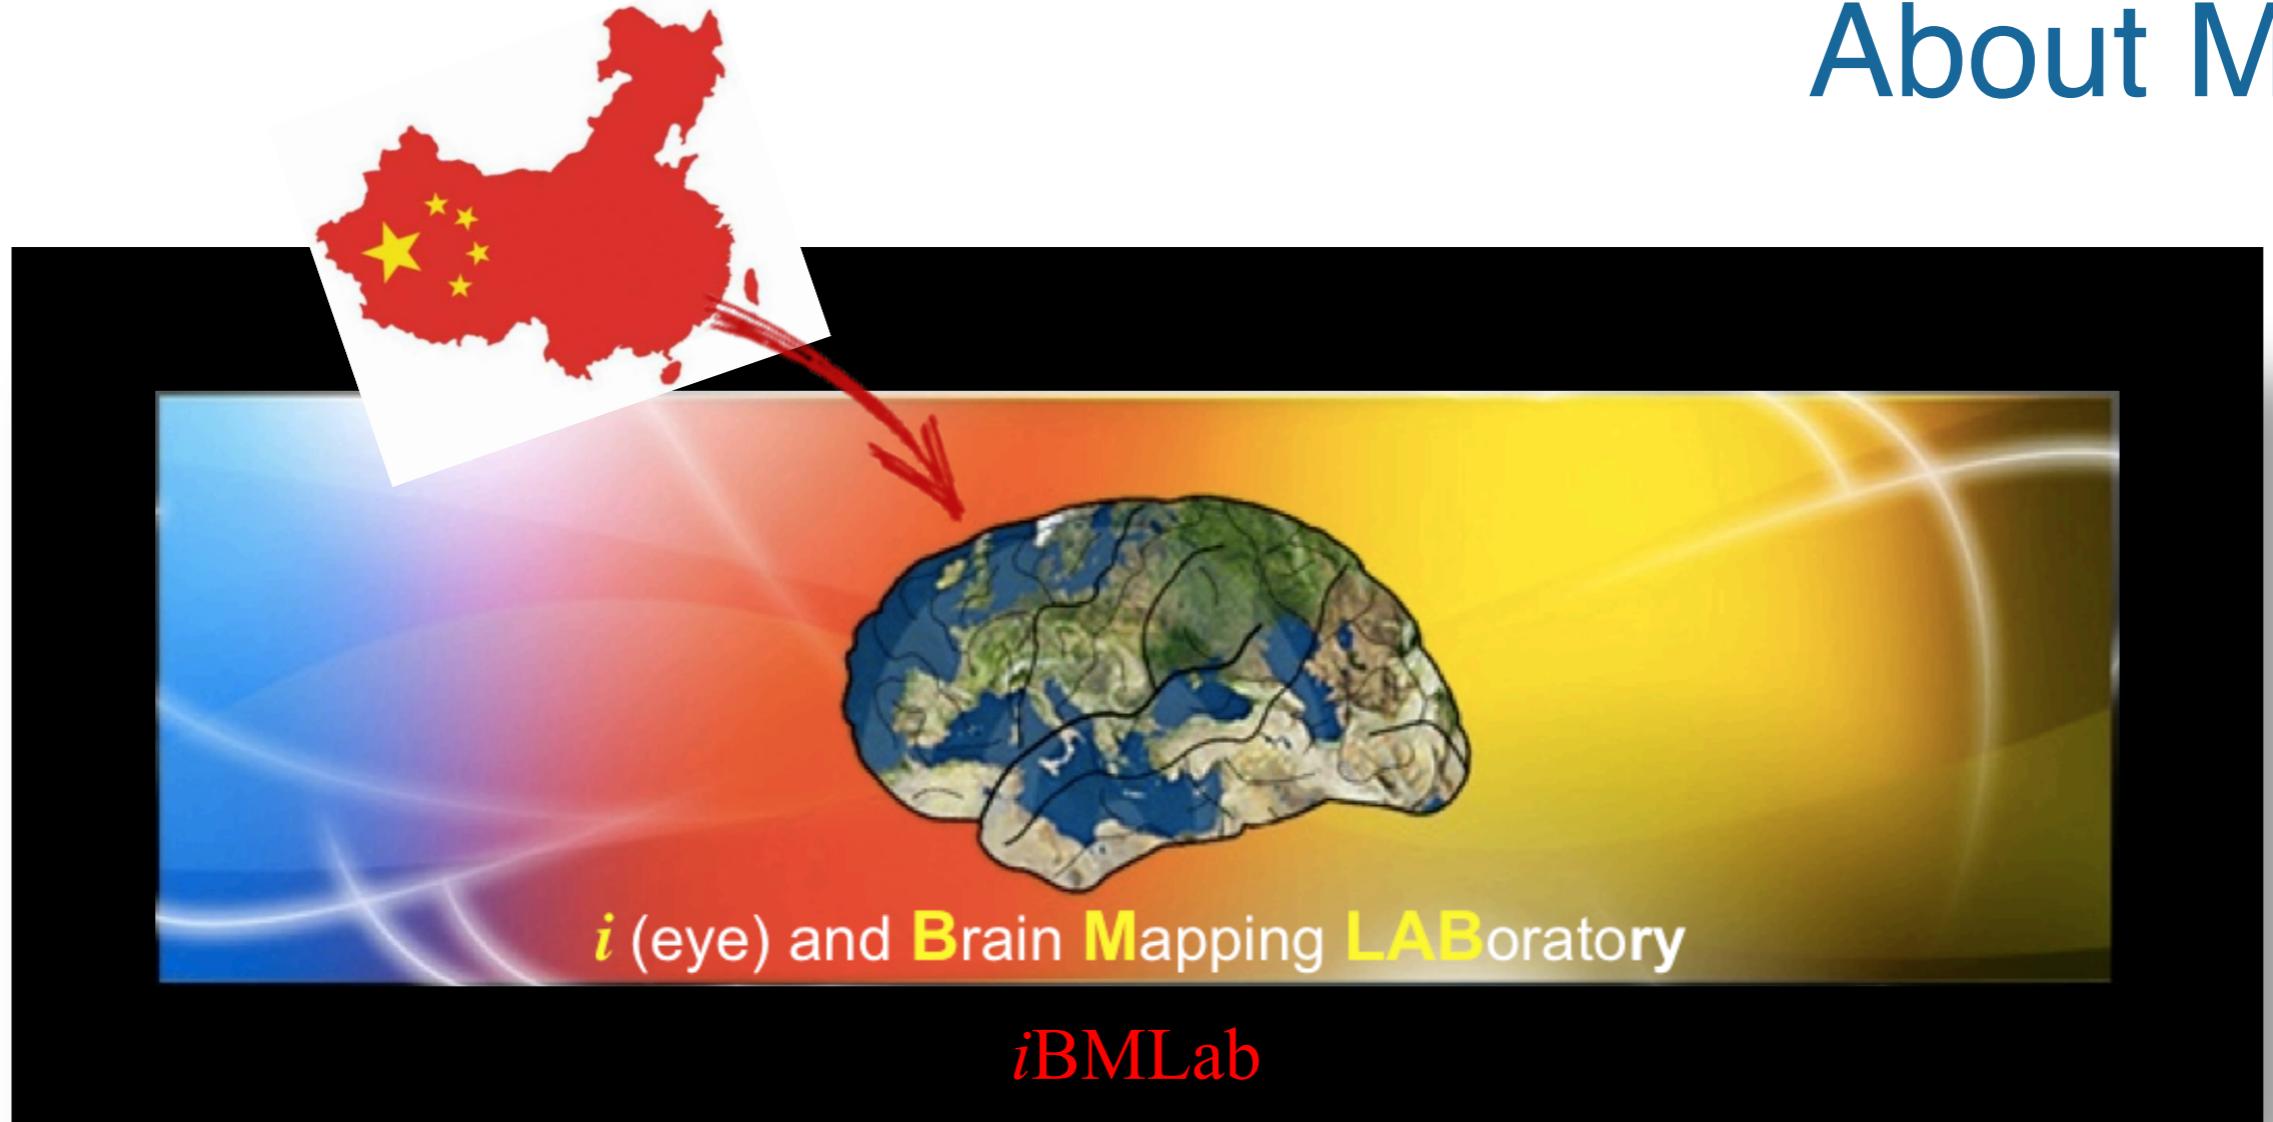
Junpeng Lao

May 2018 @ CEAi

Powered by



About Me



University
of Glasgow

UNI
FR
■

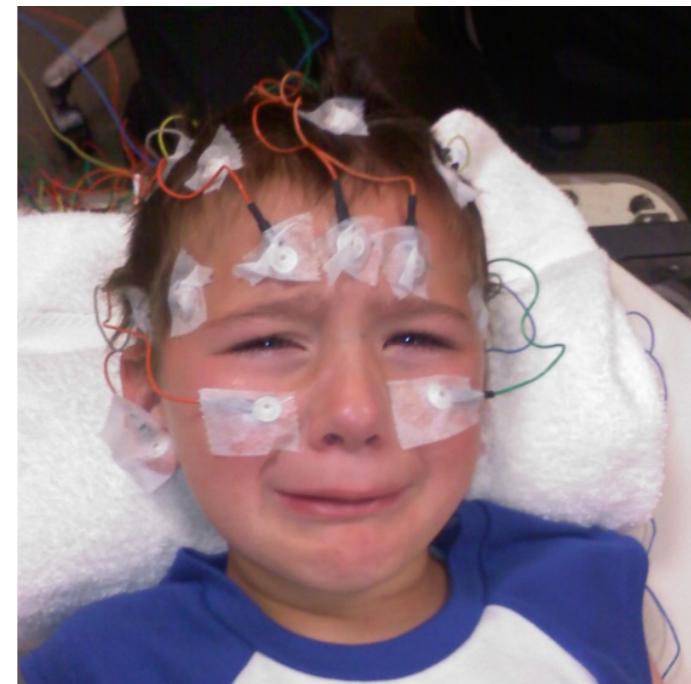
UNIVERSITÉ DE FRIBOURG
UNIVERSITÄT FREIBURG

About Me

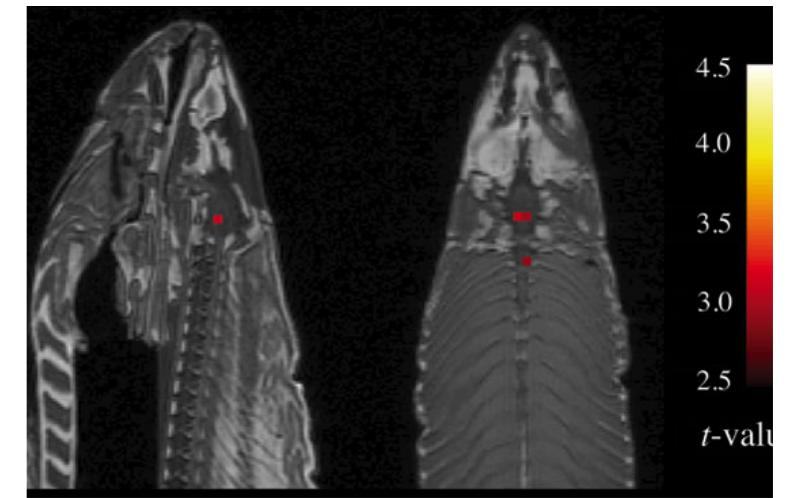
Eye-tracking



EEG



fMRI



<https://github.com/junpenglao/>
<http://junpenglao.xyz>

[all categories](#)[Latest](#)[Top](#)[Categories](#)[+ New Topic](#)

Topic	Category	Users	Replies	Views	Activity
Setting pymc3 random seed at once	Questions	2 users	1	14	2h
<input checked="" type="checkbox"/> Using WAIC to compare Log Predictive Accuracy	Questions	3 users	3	6	2h
Concepts of Parameter Estimation and Predictions, and Out of Sample Predicted Probability for Logistic Regression	Questions	2 users	5	36	4h
<input checked="" type="checkbox"/> ValueError: Bad initial energy: inf. The model might be misspecified error replicating Multilevel Regression and Poststratification with PyMC3 notebook	Questions	2 users	2	13	4h
Metropolis: ideal balance between slow-mixing and high rejection rate?	Questions	3 users	3	13	5h
<input checked="" type="checkbox"/> The pymc3 way: Linear regression and inferring given posterior/trace	Questions	3 users	4	26	18h
<input checked="" type="checkbox"/> Autocorrelation in a model	Questions	2 users	8	35	22h

<https://discourse.pymc.io/>



Prerequisites:

Branch: master ▾	PrecisionWorkshop1_Prep / notebooks /	Create new file	Upload files	Find file	History
	This branch is 131 commits ahead of fonnebeck:master.				
	 jurah Merge remote-tracking branch 'origin/master' into typos-201804				Latest commit 300113e 12 hours ago
	..				
	 images committed face model diagrams for notebook 2a				2 months ago
	 0. Introduction to the material.ipynb minor reformulation changes to introductory notebooks				a month ago
	 1a. Introduction to PyMC3.ipynb fix some typos in text of 1a				22 days ago
	 1b. Case study in PyMC3 - Coal m... overhauled initial notebooks 0 and 1				2 months ago
	 1c. Analytical solution: Beta Binom... overhauled initial notebooks 0 and 1				2 months ago
	 1d. Analytical solution: Normal-No... overhauled initial notebooks 0 and 1				2 months ago
	 2a. Model Building with PyMC3.ip... added joint plot of eye positions to make it clear how degenerate sol...				a month ago
	 2b. Tagging name tokens.ipynb removed trueskill notebooks and split face/tagging case study				2 months ago
	 3a. Sampling algorithms for Bayes... Add reference to 3a notebook to add more info about law of large numbers				a month ago
	 3b. Importance sampling.ipynb added Cauchy proposal distribution to 3b				a month ago
	 3c. Rejection sampling.ipynb incorporated feedback into notebooks 3b-3c				a month ago
	 3d. MC in more than one dimensio... fix some typos				a month ago
	 3e. Introduction to MCMC.ipynb typos in text of 3e				22 days ago
	 3f. The Metropolis sampler.ipynb typo in 3f				21 days ago
	 3g. Case study: reaction times of ... typos in 3g				22 days ago
	 3g.1 Case study: reaction times of... added solution to notebook 3g as 3g.1, added more text to notebook 3i				22 days ago



Objective of the workshop:

Understand what is being computed Bayesian modelling

Learn about ways to validate your model in PyMC3

Learn about different ways to perform proper inference

Evaluation of model fitting and model comparisons



Probability theory

Measure-Theoretical Foundations of Probability Theory

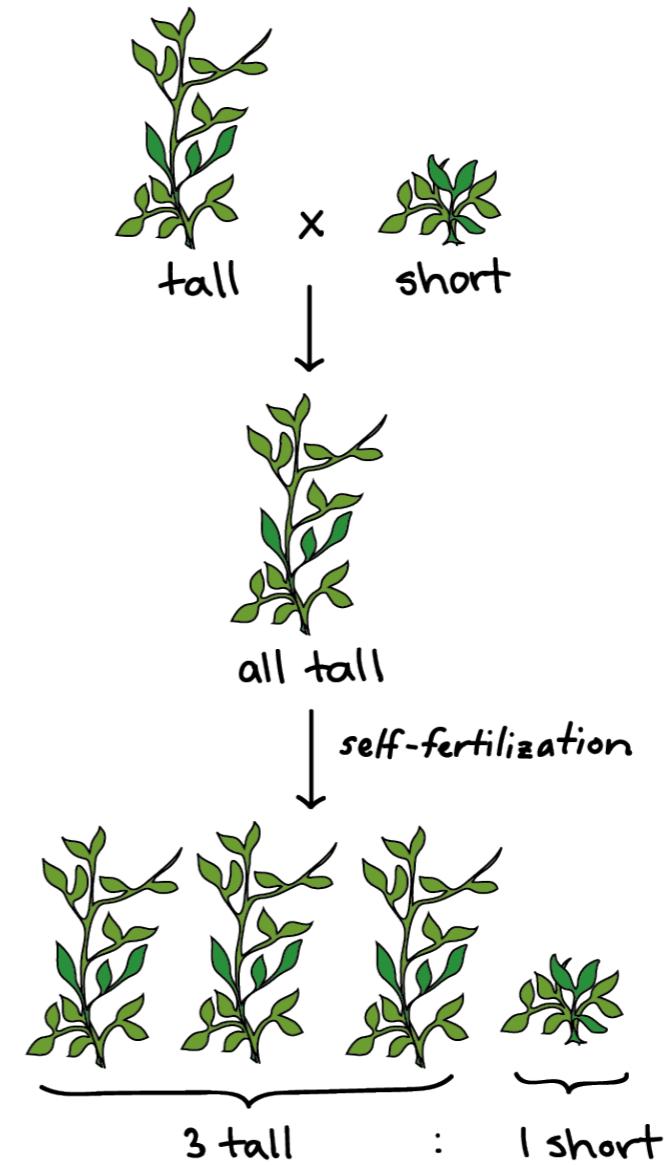
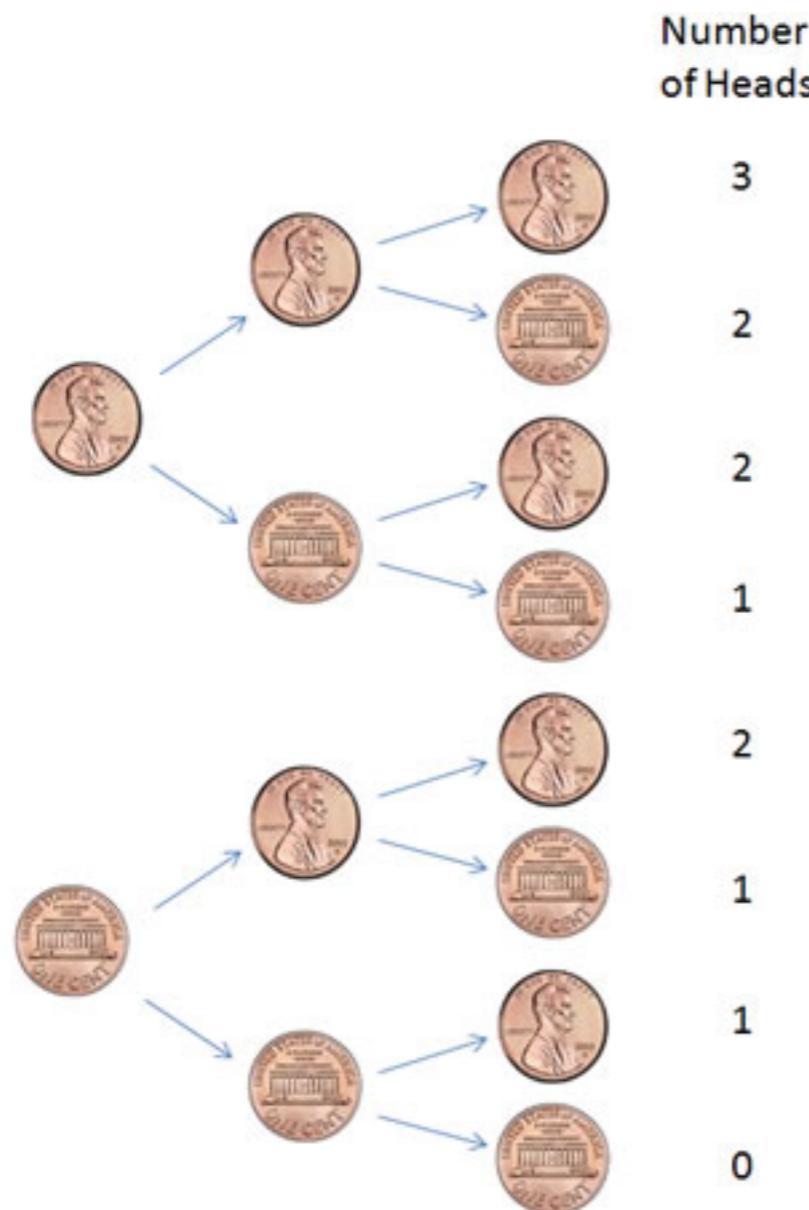
- Measure theory
 - σ -algebra
- Measurable mapping
 - pushforward measure
- Integral
 - expectation



https://betanalpha.github.io/assets/case_studies/probability_theory.html



Random Variable

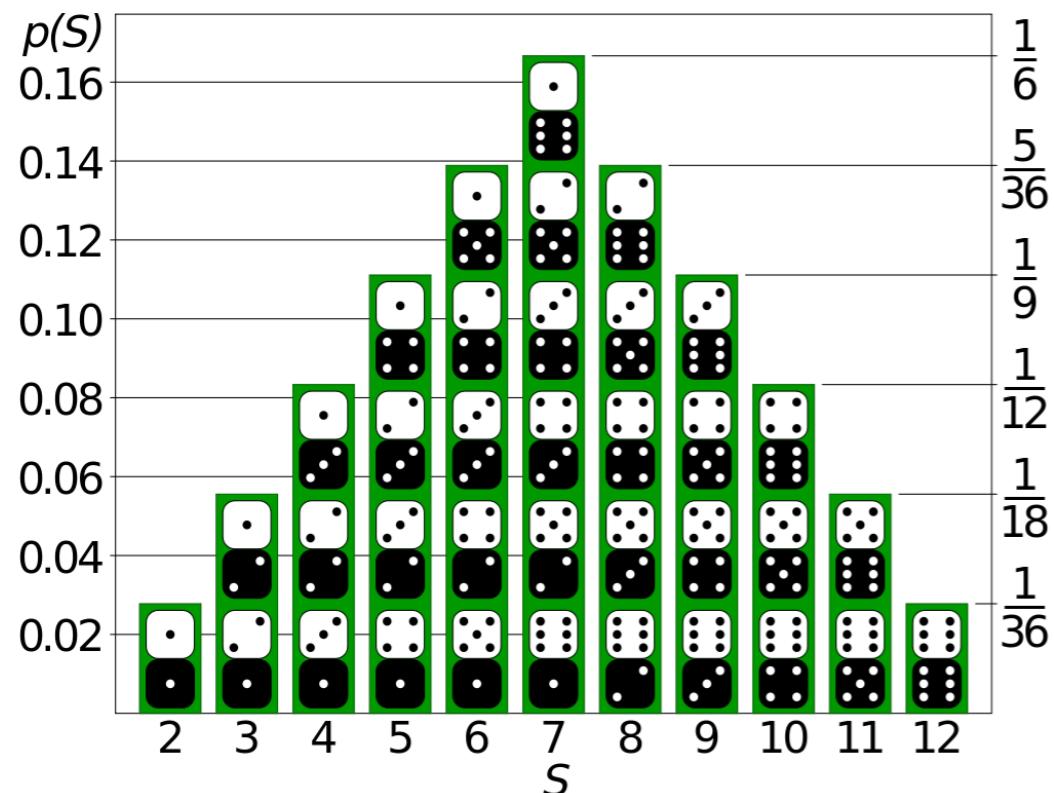


Mendel's actual numbers:
787 tall : 277 short (2.84:1).

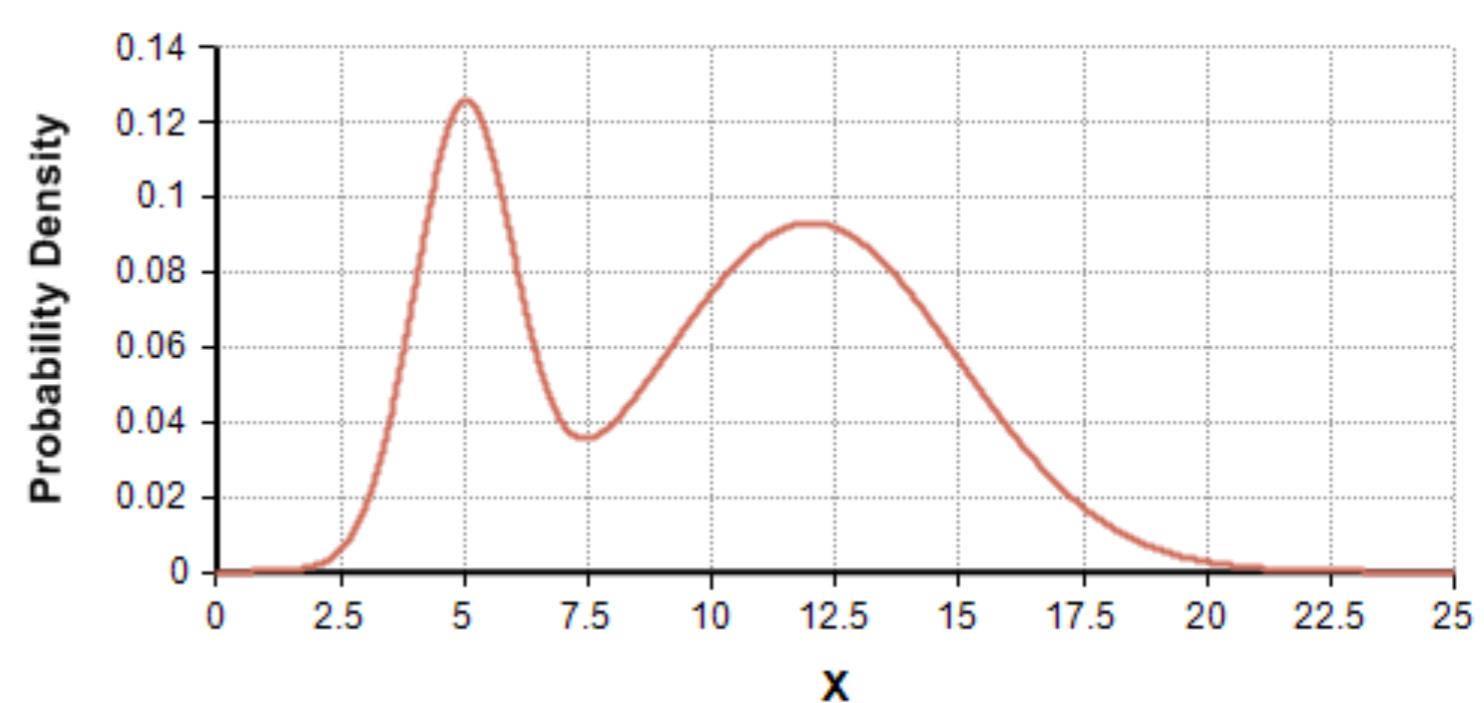


Probability distribution

Discrete case:



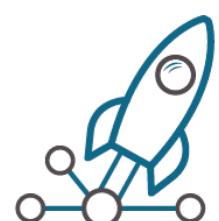
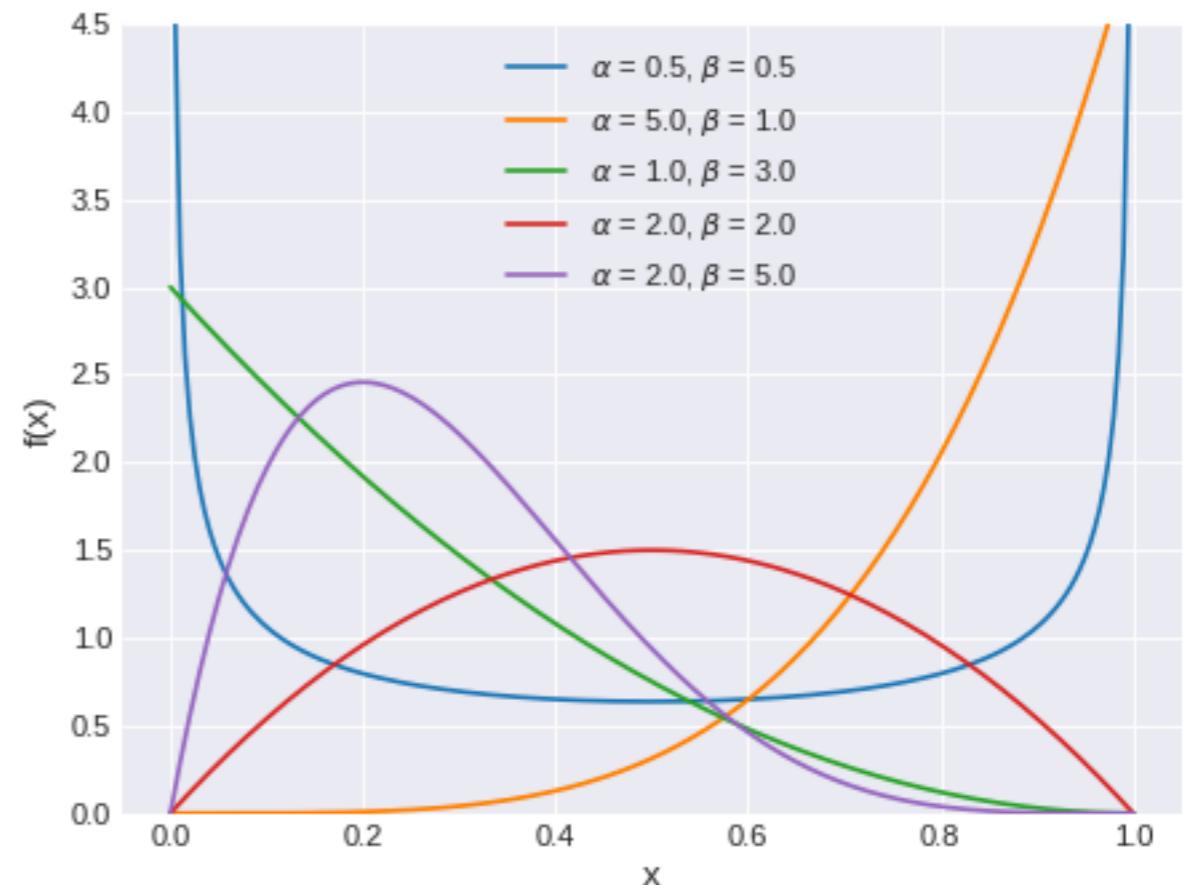
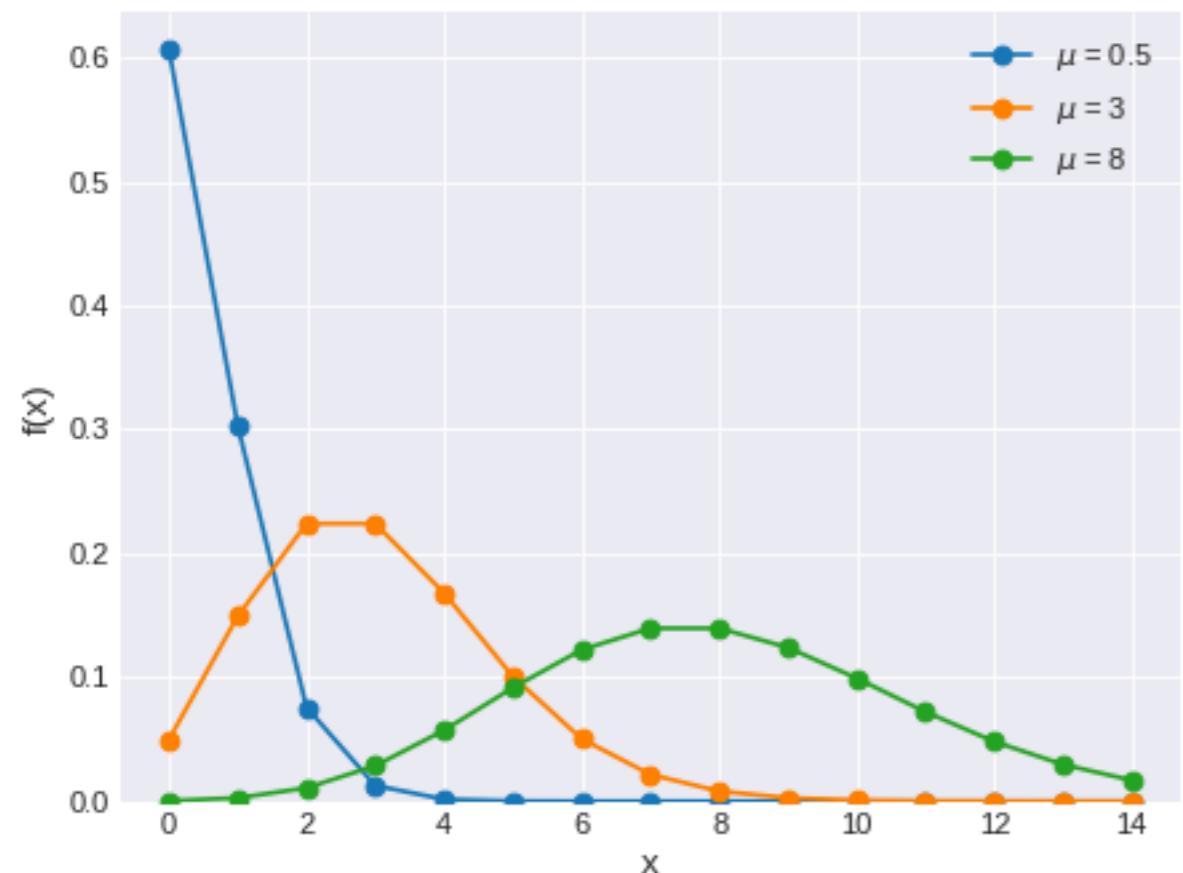
Continuous case



Probability distribution

$$f(x | \mu) = \frac{e^{-\mu} \mu^x}{x!}$$

$$f(x | \alpha, \beta) = \frac{x^{\alpha-1} (1-x)^{\beta-1}}{B(\alpha, \beta)}$$



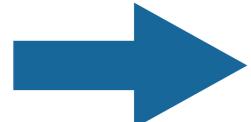
Random Variable

$$y \sim \pi(\theta)$$

“Bob is flipping a fair coin and he stop when he observed 5 heads”

X = the number of total coin flips

```
p = .5
flip_sq = st.bernoulli.rvs(p, size=5) # first 5 flips
while flip_sq.sum() != 5:
    flip_sq = np.append(flip_sq, st.bernoulli.rvs(p))
print(flip_sq)
```



Code1 - Forward_Random.ipynb



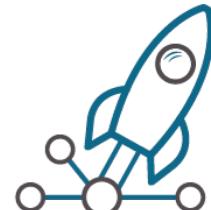
Random Variable

“Bob is flipping a fair coin and he stop when he observed 5 heads”

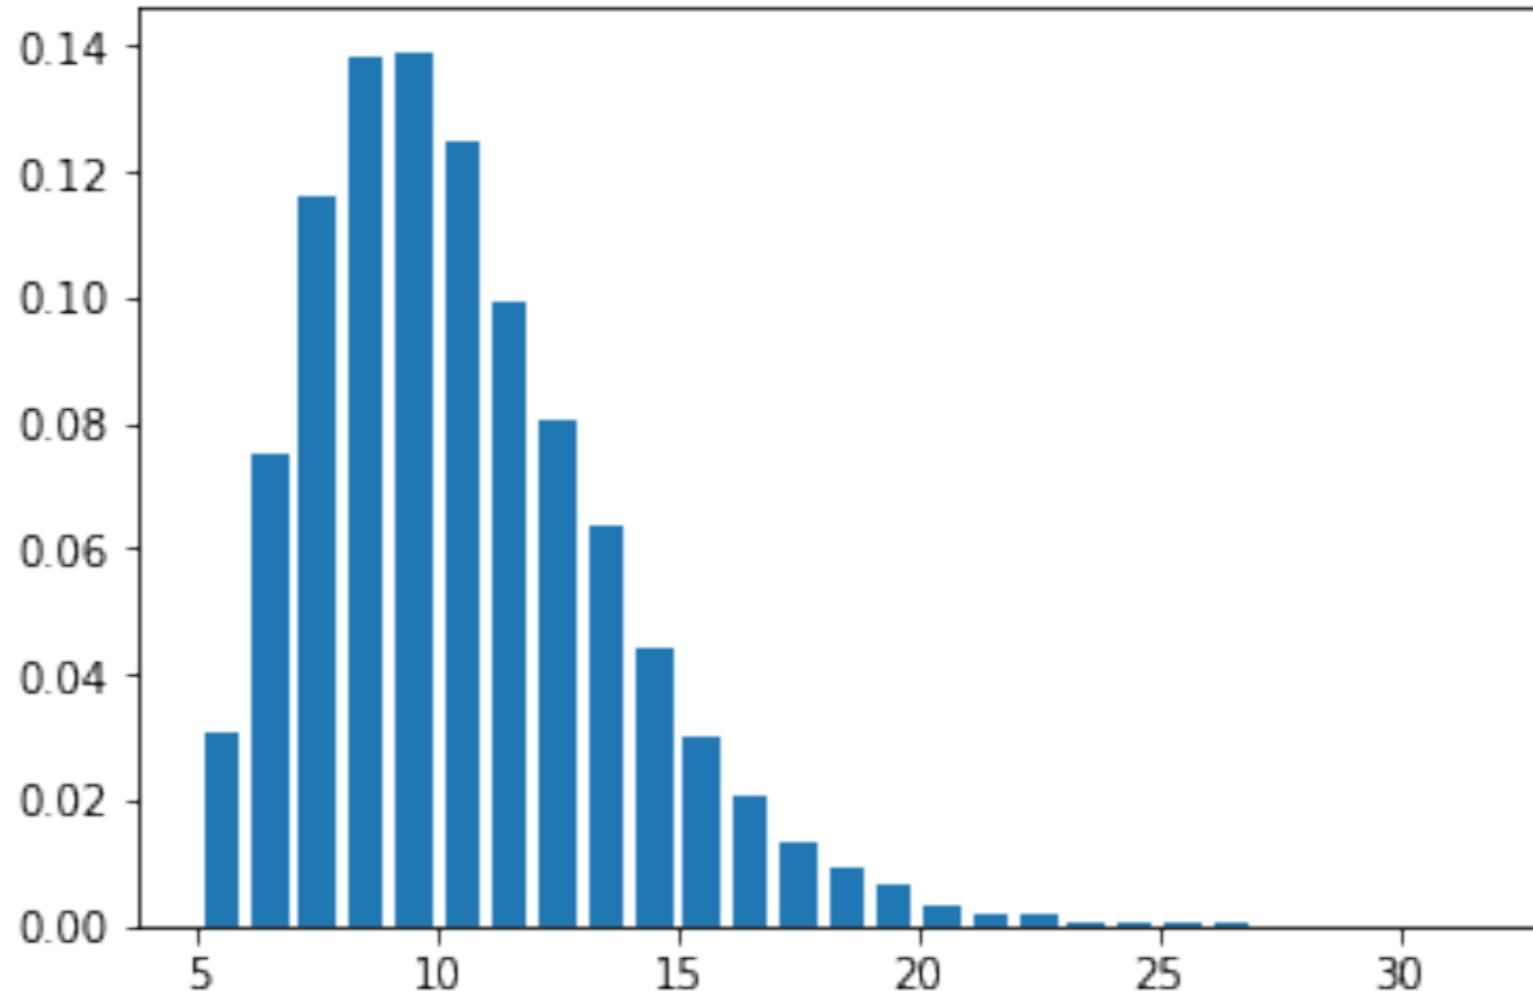
```
p = .5
flip_sq = st.bernoulli.rvs(p, size=5) # first 5 flips
while flip_sq.sum() != 5:
    flip_sq = np.append(flip_sq, st.bernoulli.rvs(p))
print(flip_sq)
```

X = the number of total coin flips

```
7 <-- [0 1 1 1 0 1 1]
9 <-- [1 0 1 1 0 0 0 1 1]
10 <-- [0 0 1 0 1 0 0 1 1 1]
9 <-- [1 0 0 0 0 1 1 1 1]
11 <-- [0 1 1 1 0 0 1 0 0 0 1]
10 <-- [1 0 0 1 0 0 1 1 0 1]
6 <-- [0 1 1 1 1 1]
15 <-- [0 0 0 1 0 1 0 1 0 0 0 0 1 0 1]
11 <-- [0 0 0 1 0 1 1 1 0 0 1]
11 <-- [1 0 0 0 0 1 1 1 0 0 1]
```



Underlying distribution

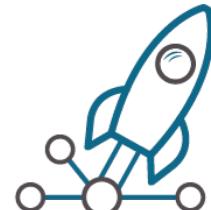
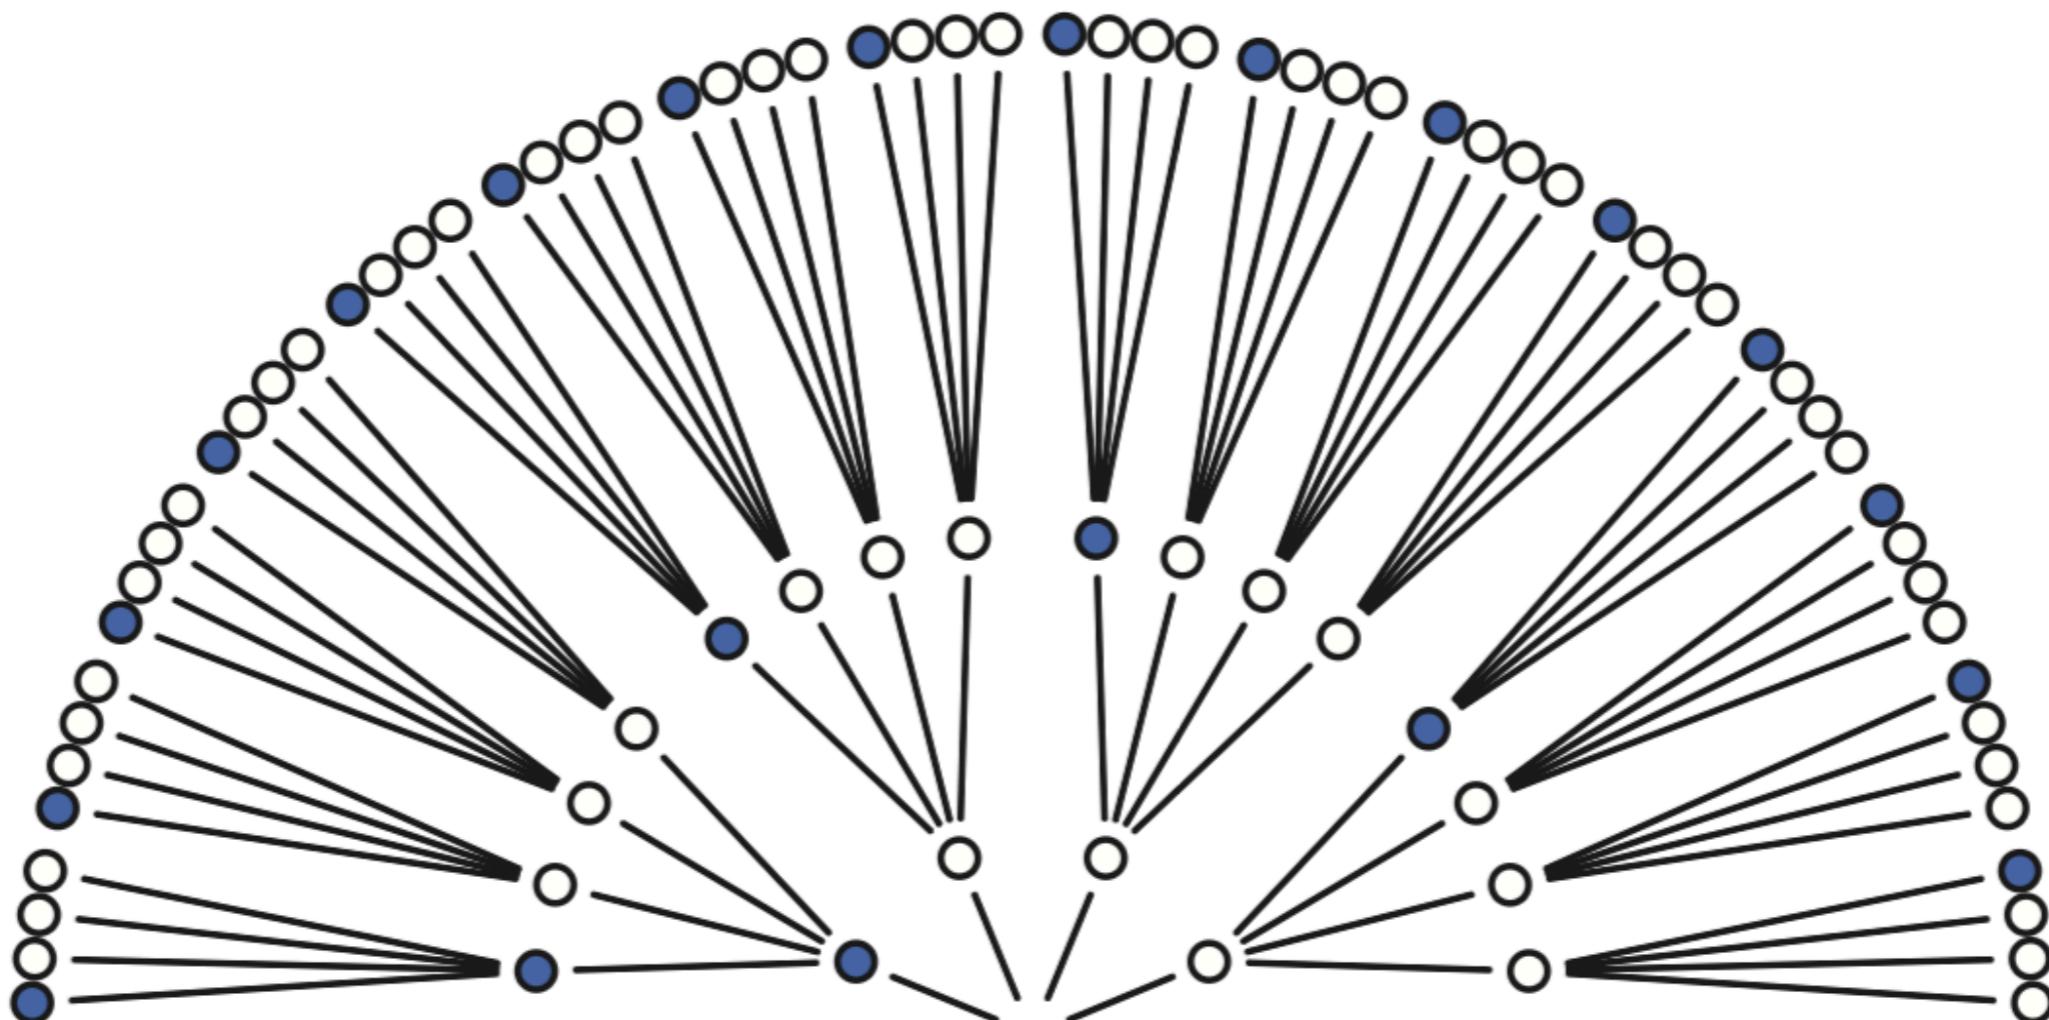


Negative binomial(nlp, y)



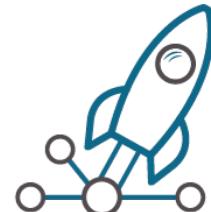
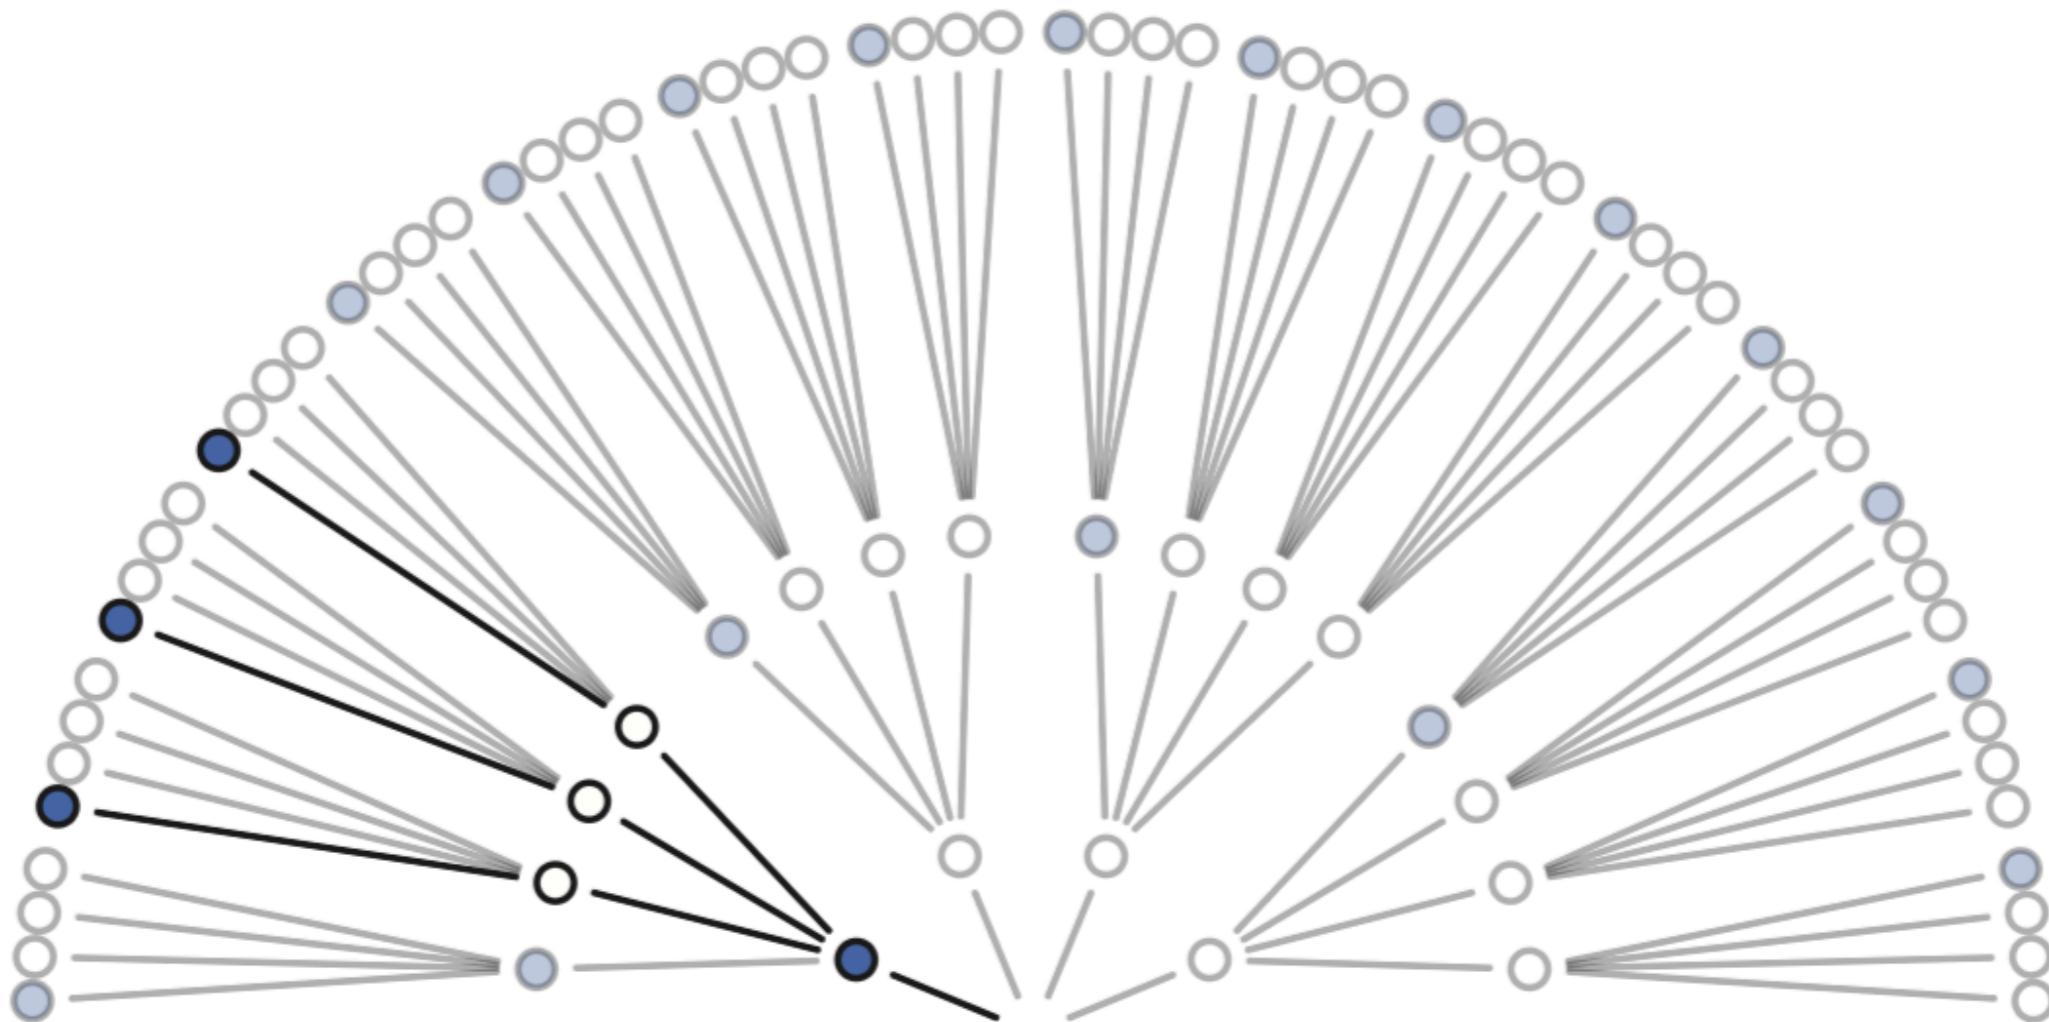
The garden of forking data

From the 2nd chapter of Richard McElreath's **Statistical Rethinking**



The garden of forking data

Observed: ● ○ ●



The garden of forking data

Observed: ● ○ ●

[○○○○]

[●○○○]

[●●○○]

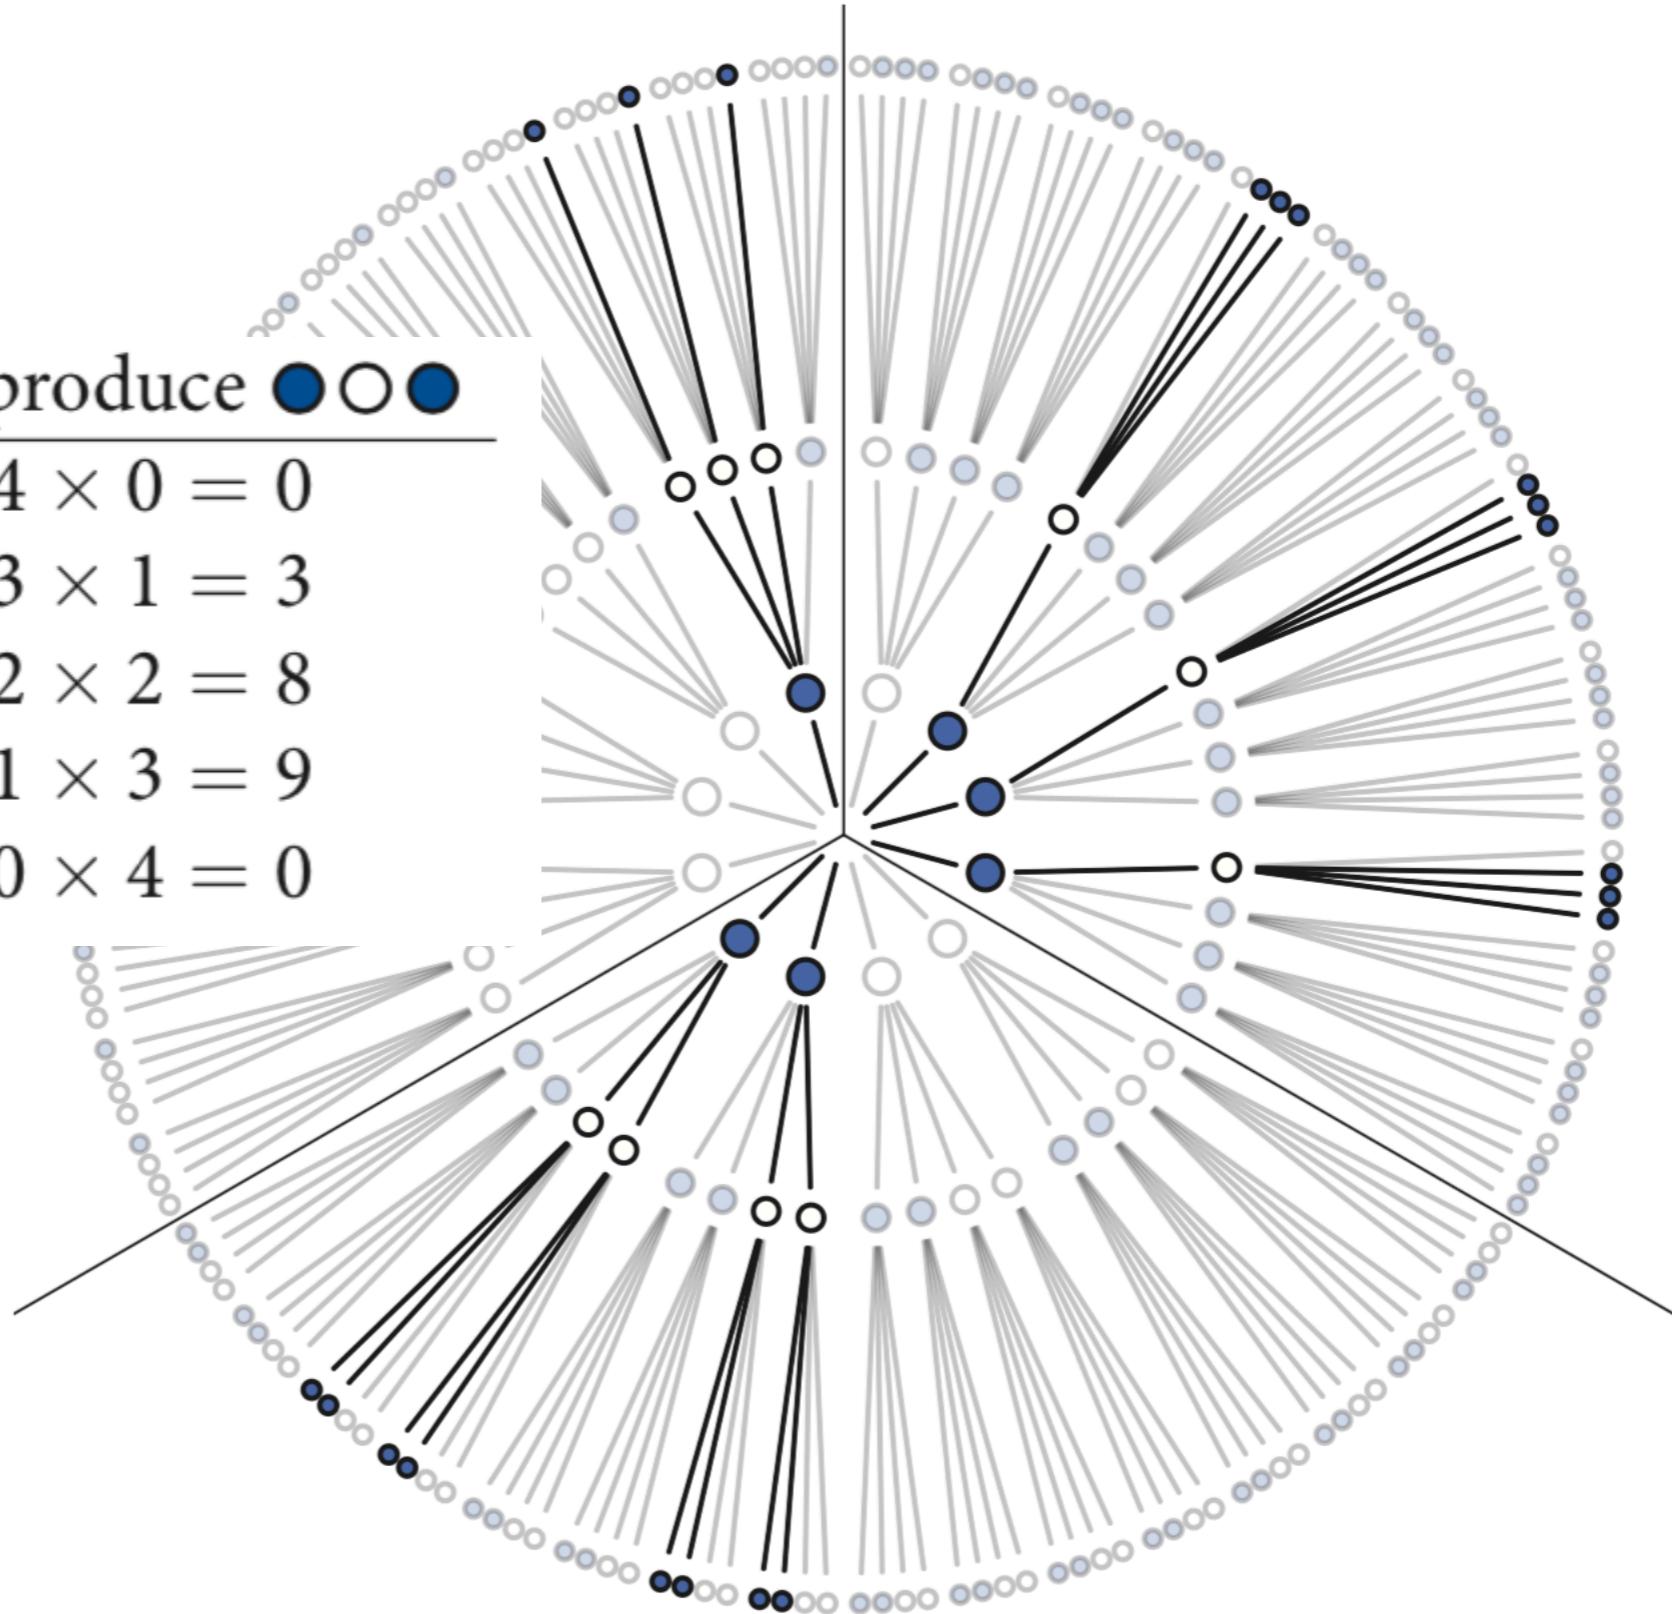
[●●●○]

[●●●●]



The garden of forking data

Conjecture	Ways to produce 
[○○○○]	$0 \times 4 \times 0 = 0$
[●○○○]	$1 \times 3 \times 1 = 3$
[●●○○]	$2 \times 2 \times 2 = 8$
[●●●○]	$3 \times 1 \times 3 = 9$
[●●●●]	$4 \times 0 \times 4 = 0$



Counting the paths within a small world

“Bob is flipping a fair coin and he stop when he observed 5 heads”

[1 1 1 1 1]

[0 1 1 1 1 1] * 5

...

“He ends up flipping the coin 9 times”

“The bias of the coin is either 0.5 or 0.75”

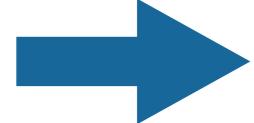


Likelihood function

For example, binomial PMF

$$\pi(x \mid n, p) = \binom{n}{x} p^x (1 - p)^{n-x}$$

$$f(x, n, p) = \binom{n}{x} p^x (1 - p)^{n-x}$$



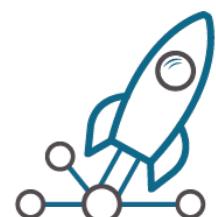
Code2 - likelihood.ipynb



Combining probabilities

Summary of probabilities

Event	Probability
A	$P(A) \in [0, 1]$
not A	$P(A^C) = 1 - P(A)$
A or B	$P(A \cup B) = P(A) + P(B) - P(A \cap B)$ $P(A \cup B) = P(A) + P(B) \quad \text{if A and B are mutually exclusive}$
A and B	$P(A \cap B) = P(A B)P(B) = P(B A)P(A)$ $P(A \cap B) = P(A)P(B) \quad \text{if A and B are independent}$
A given B	$P(A B) = \frac{P(A \cap B)}{P(B)} = \frac{P(B A)P(A)}{P(B)}$



Sum rule and product rule

Sum rule:

$$P(X) = \int_Y P(X, Y)$$

Product rule:

$$P(X, Y) = P(X \mid Y)P(Y)$$



Likelihood of observing:

$X = [0, 0, 0, 0, 1, 1]$ with $p = .7$

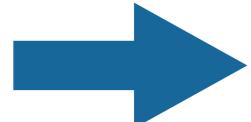
```
p = .7
y = np.asarray([0, 0, 0, 0, 1, 1])

prob = [p if y_ == 1 else (1-p) for y_ in y]
prob = np.cumprod(prob)
prob
```

```
array([0.3      , 0.09      , 0.027     , 0.0081    , 0.00567   , 0.003969])
```

```
prob2 = np.cumprod(st.bernoulli.pmf(y, p))
prob2
```

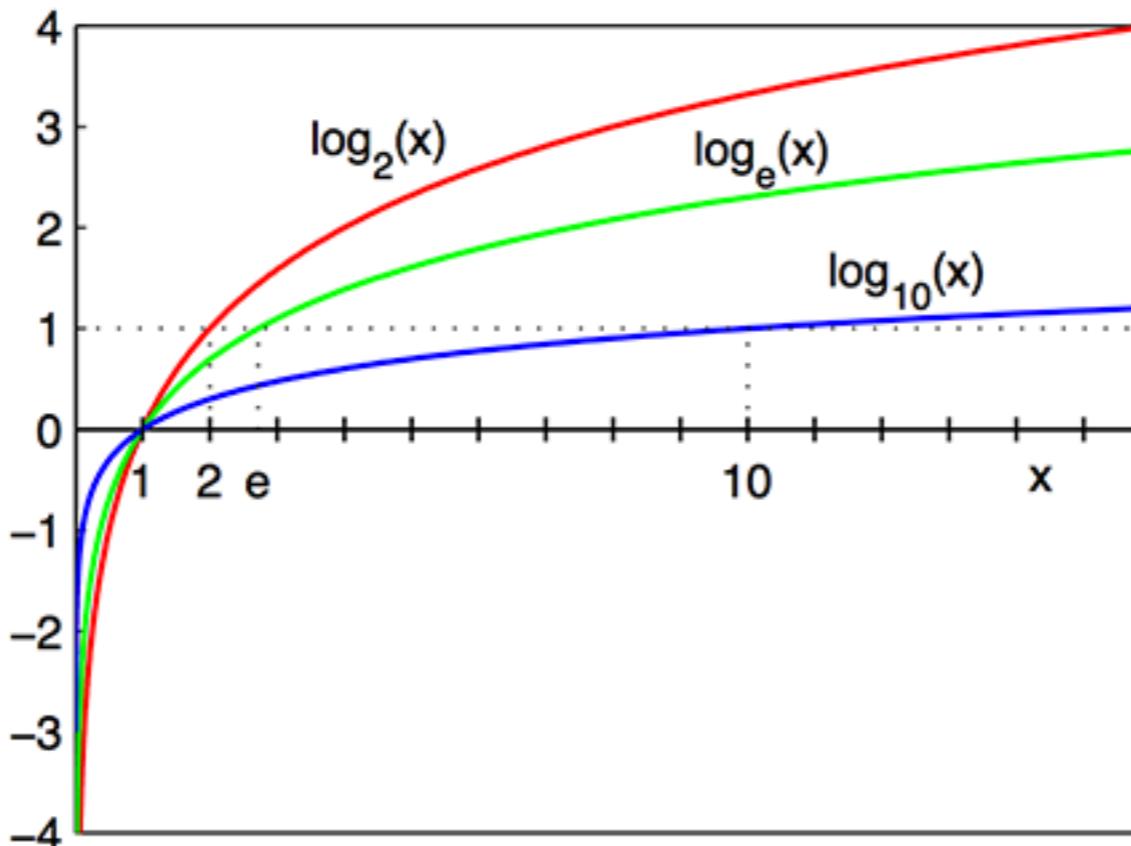
```
array([0.3      , 0.09      , 0.027     , 0.0081    , 0.00567   , 0.003969])
```



Code3 - Combining_Likelihood.ipynb



Log-Likelihood



```
np.cumsum(np.log(st.bernoulli.pmf(y, p)))
```

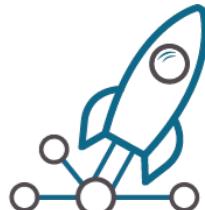
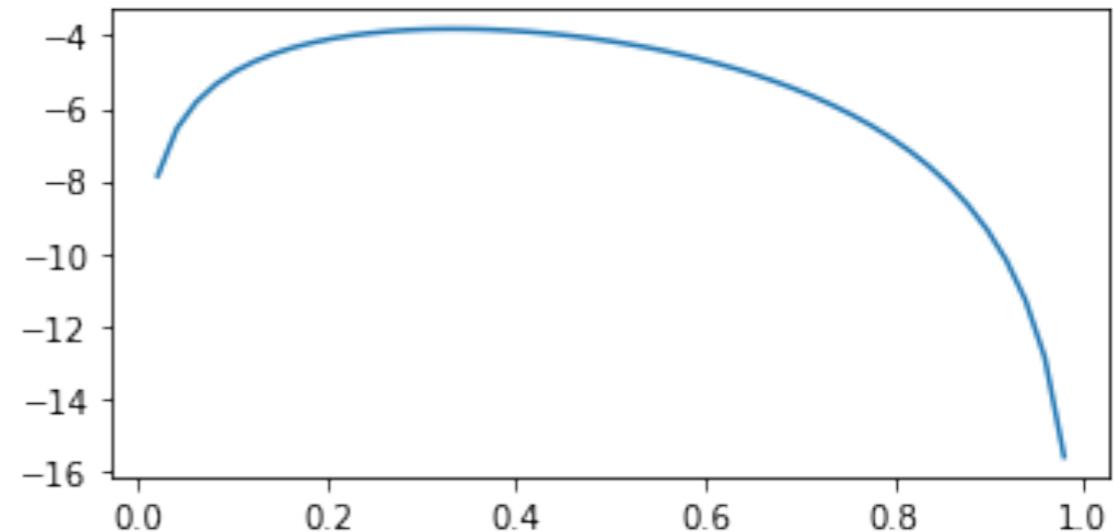
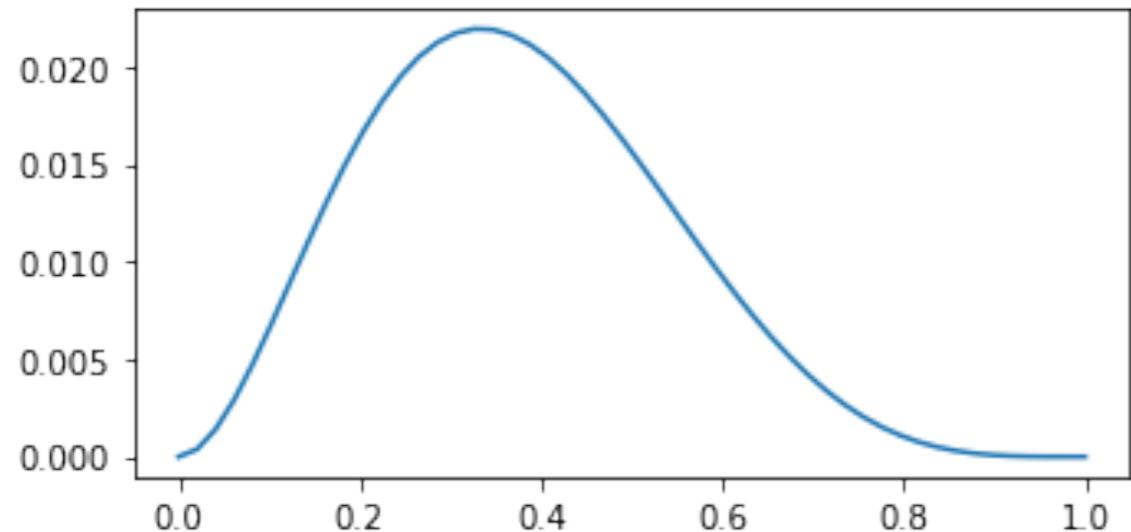
```
array([-1.2039728 , -2.40794561, -3.61191841, -4.81589122, -5.17256616,
       -5.52924111])
```

```
st.bernoulli.logpmf(y, p).cumsum()
```

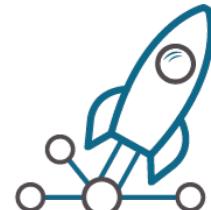
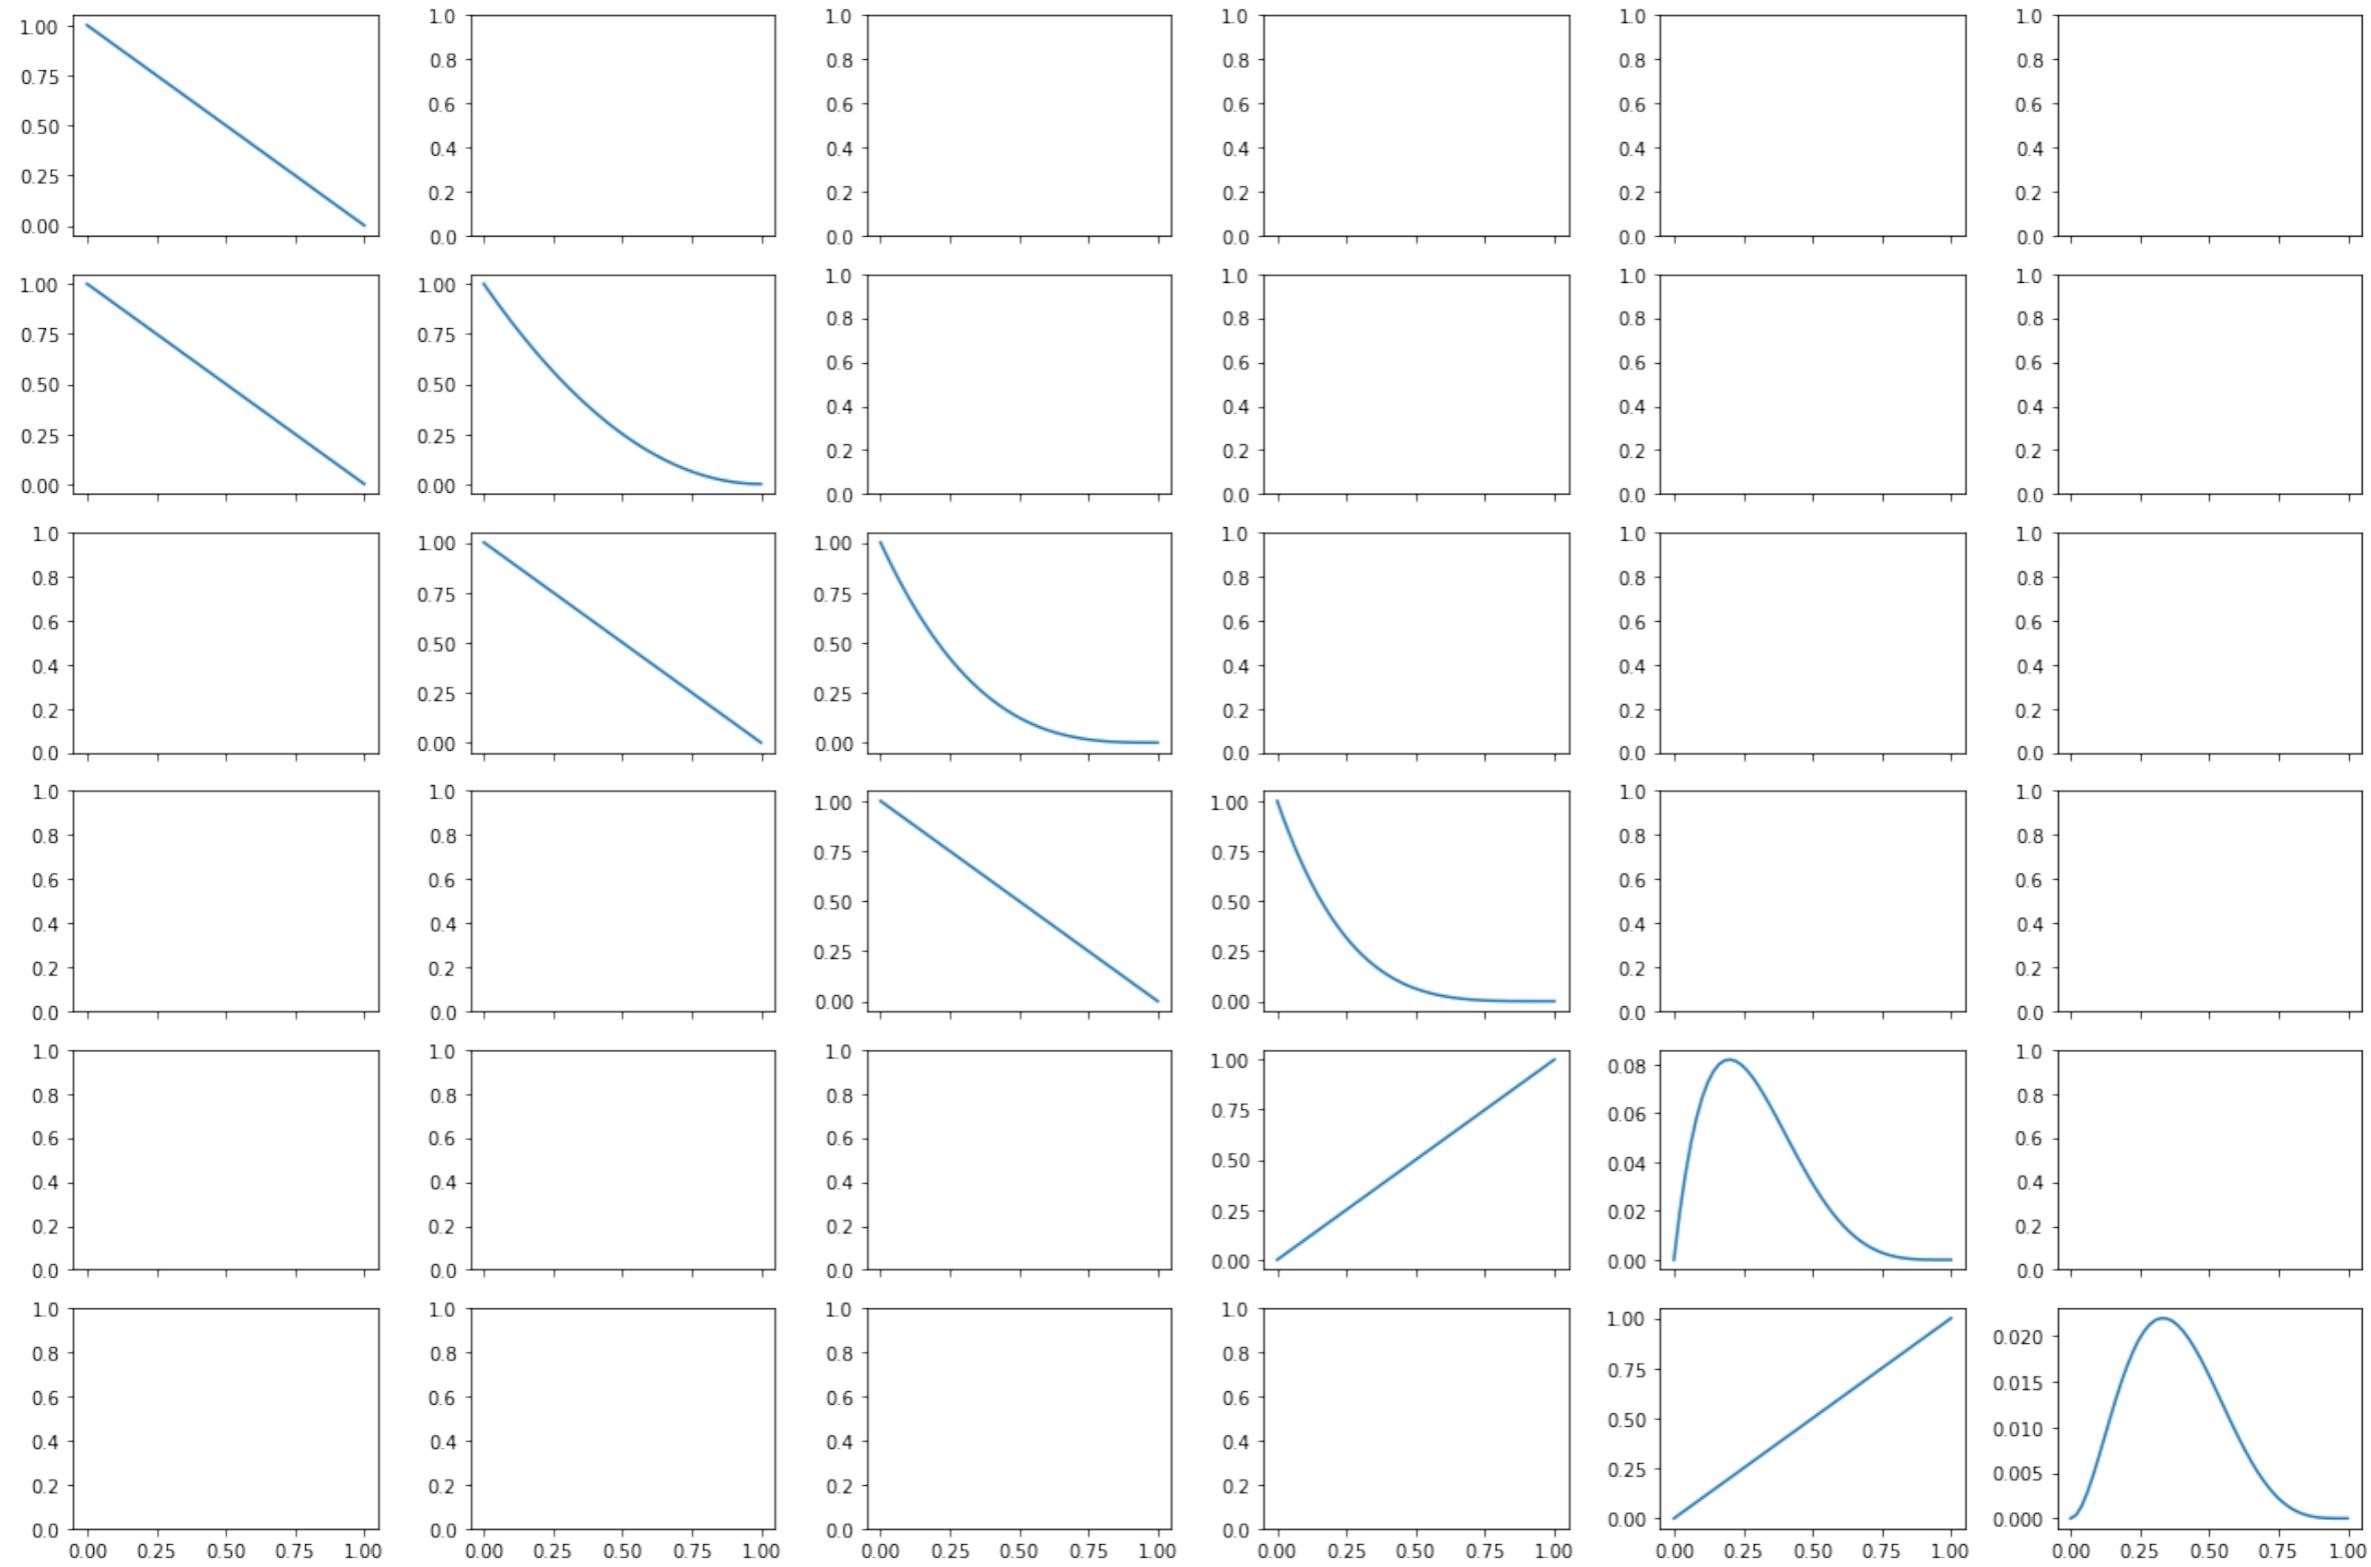


Combining Likelihood

```
lambda p: st.bernoulli.pmf(y, p).prod()
```



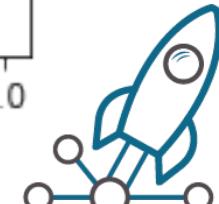
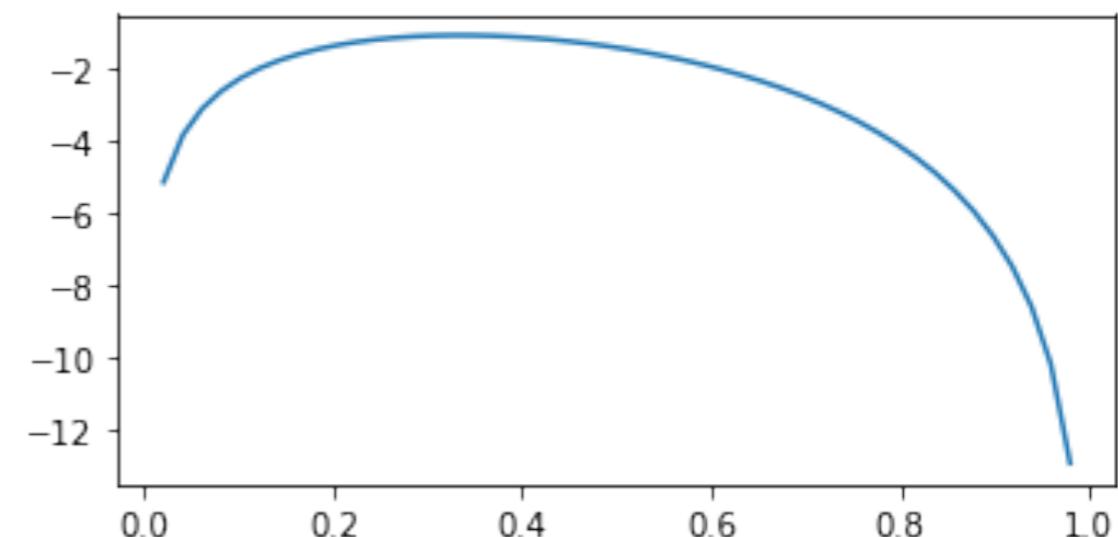
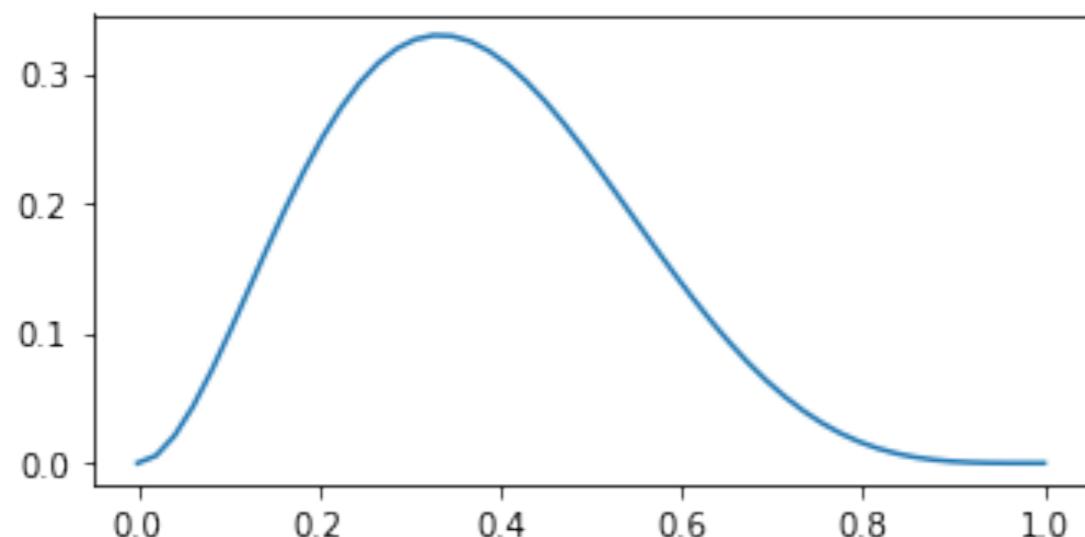
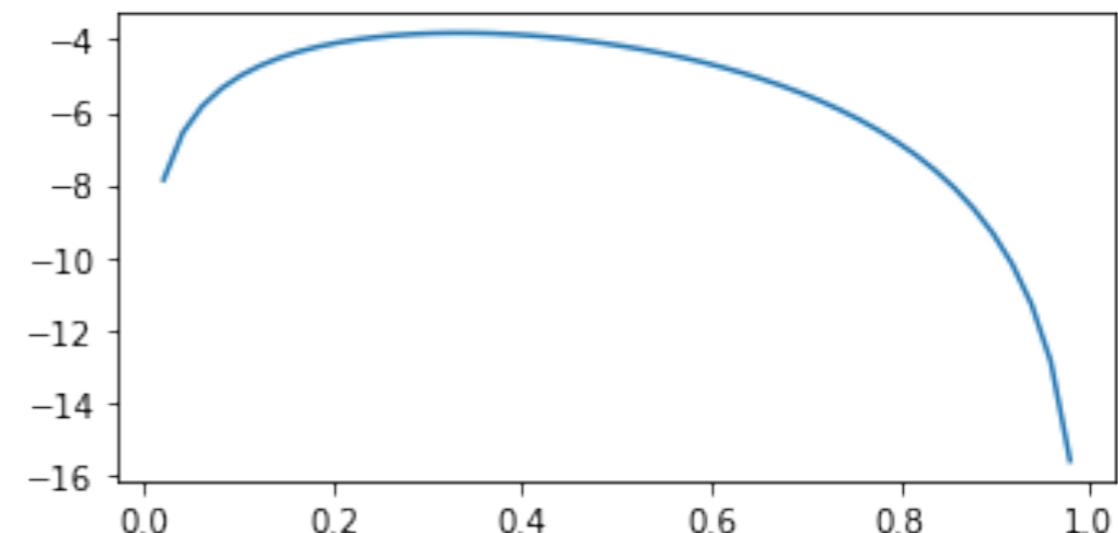
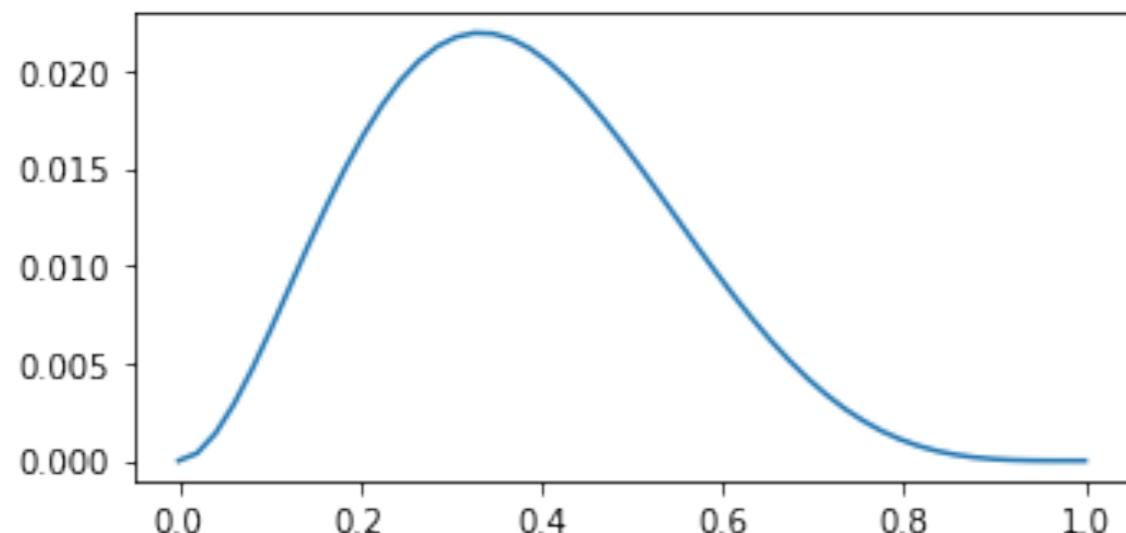
Update of Likelihood function



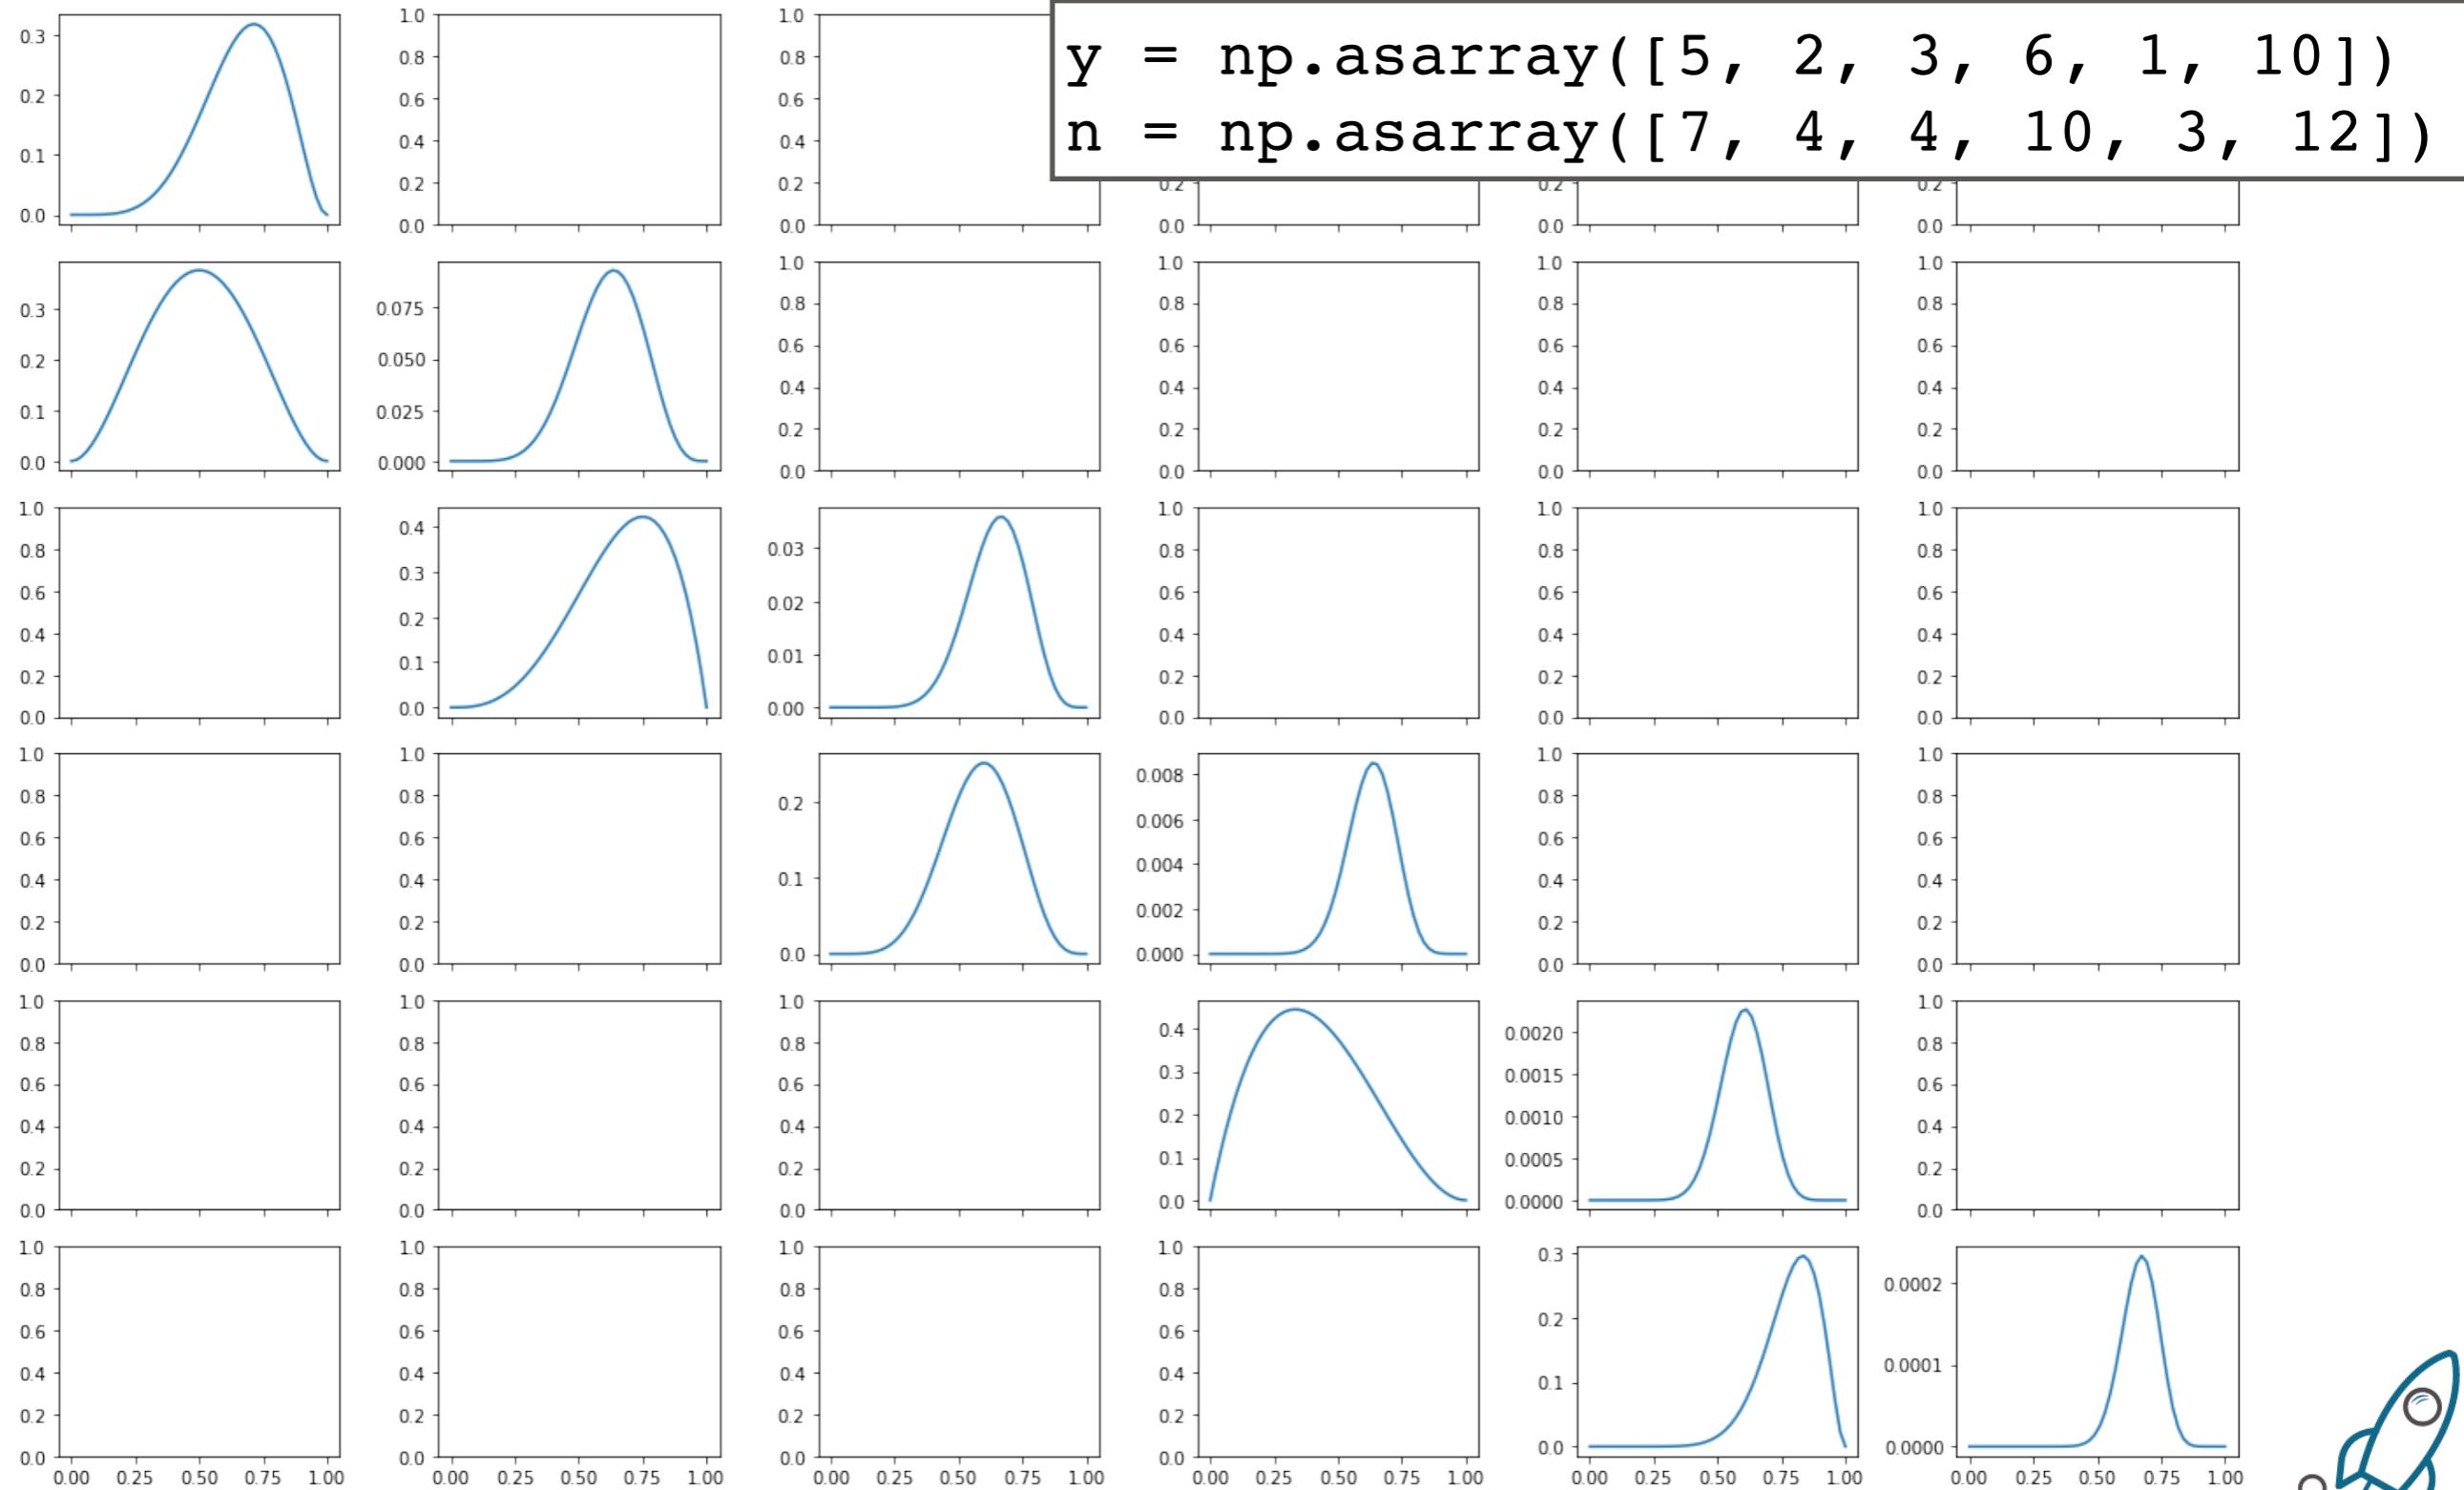
Likelihood function

lambda p: st.bernoulli.pmf(y, p).prod()

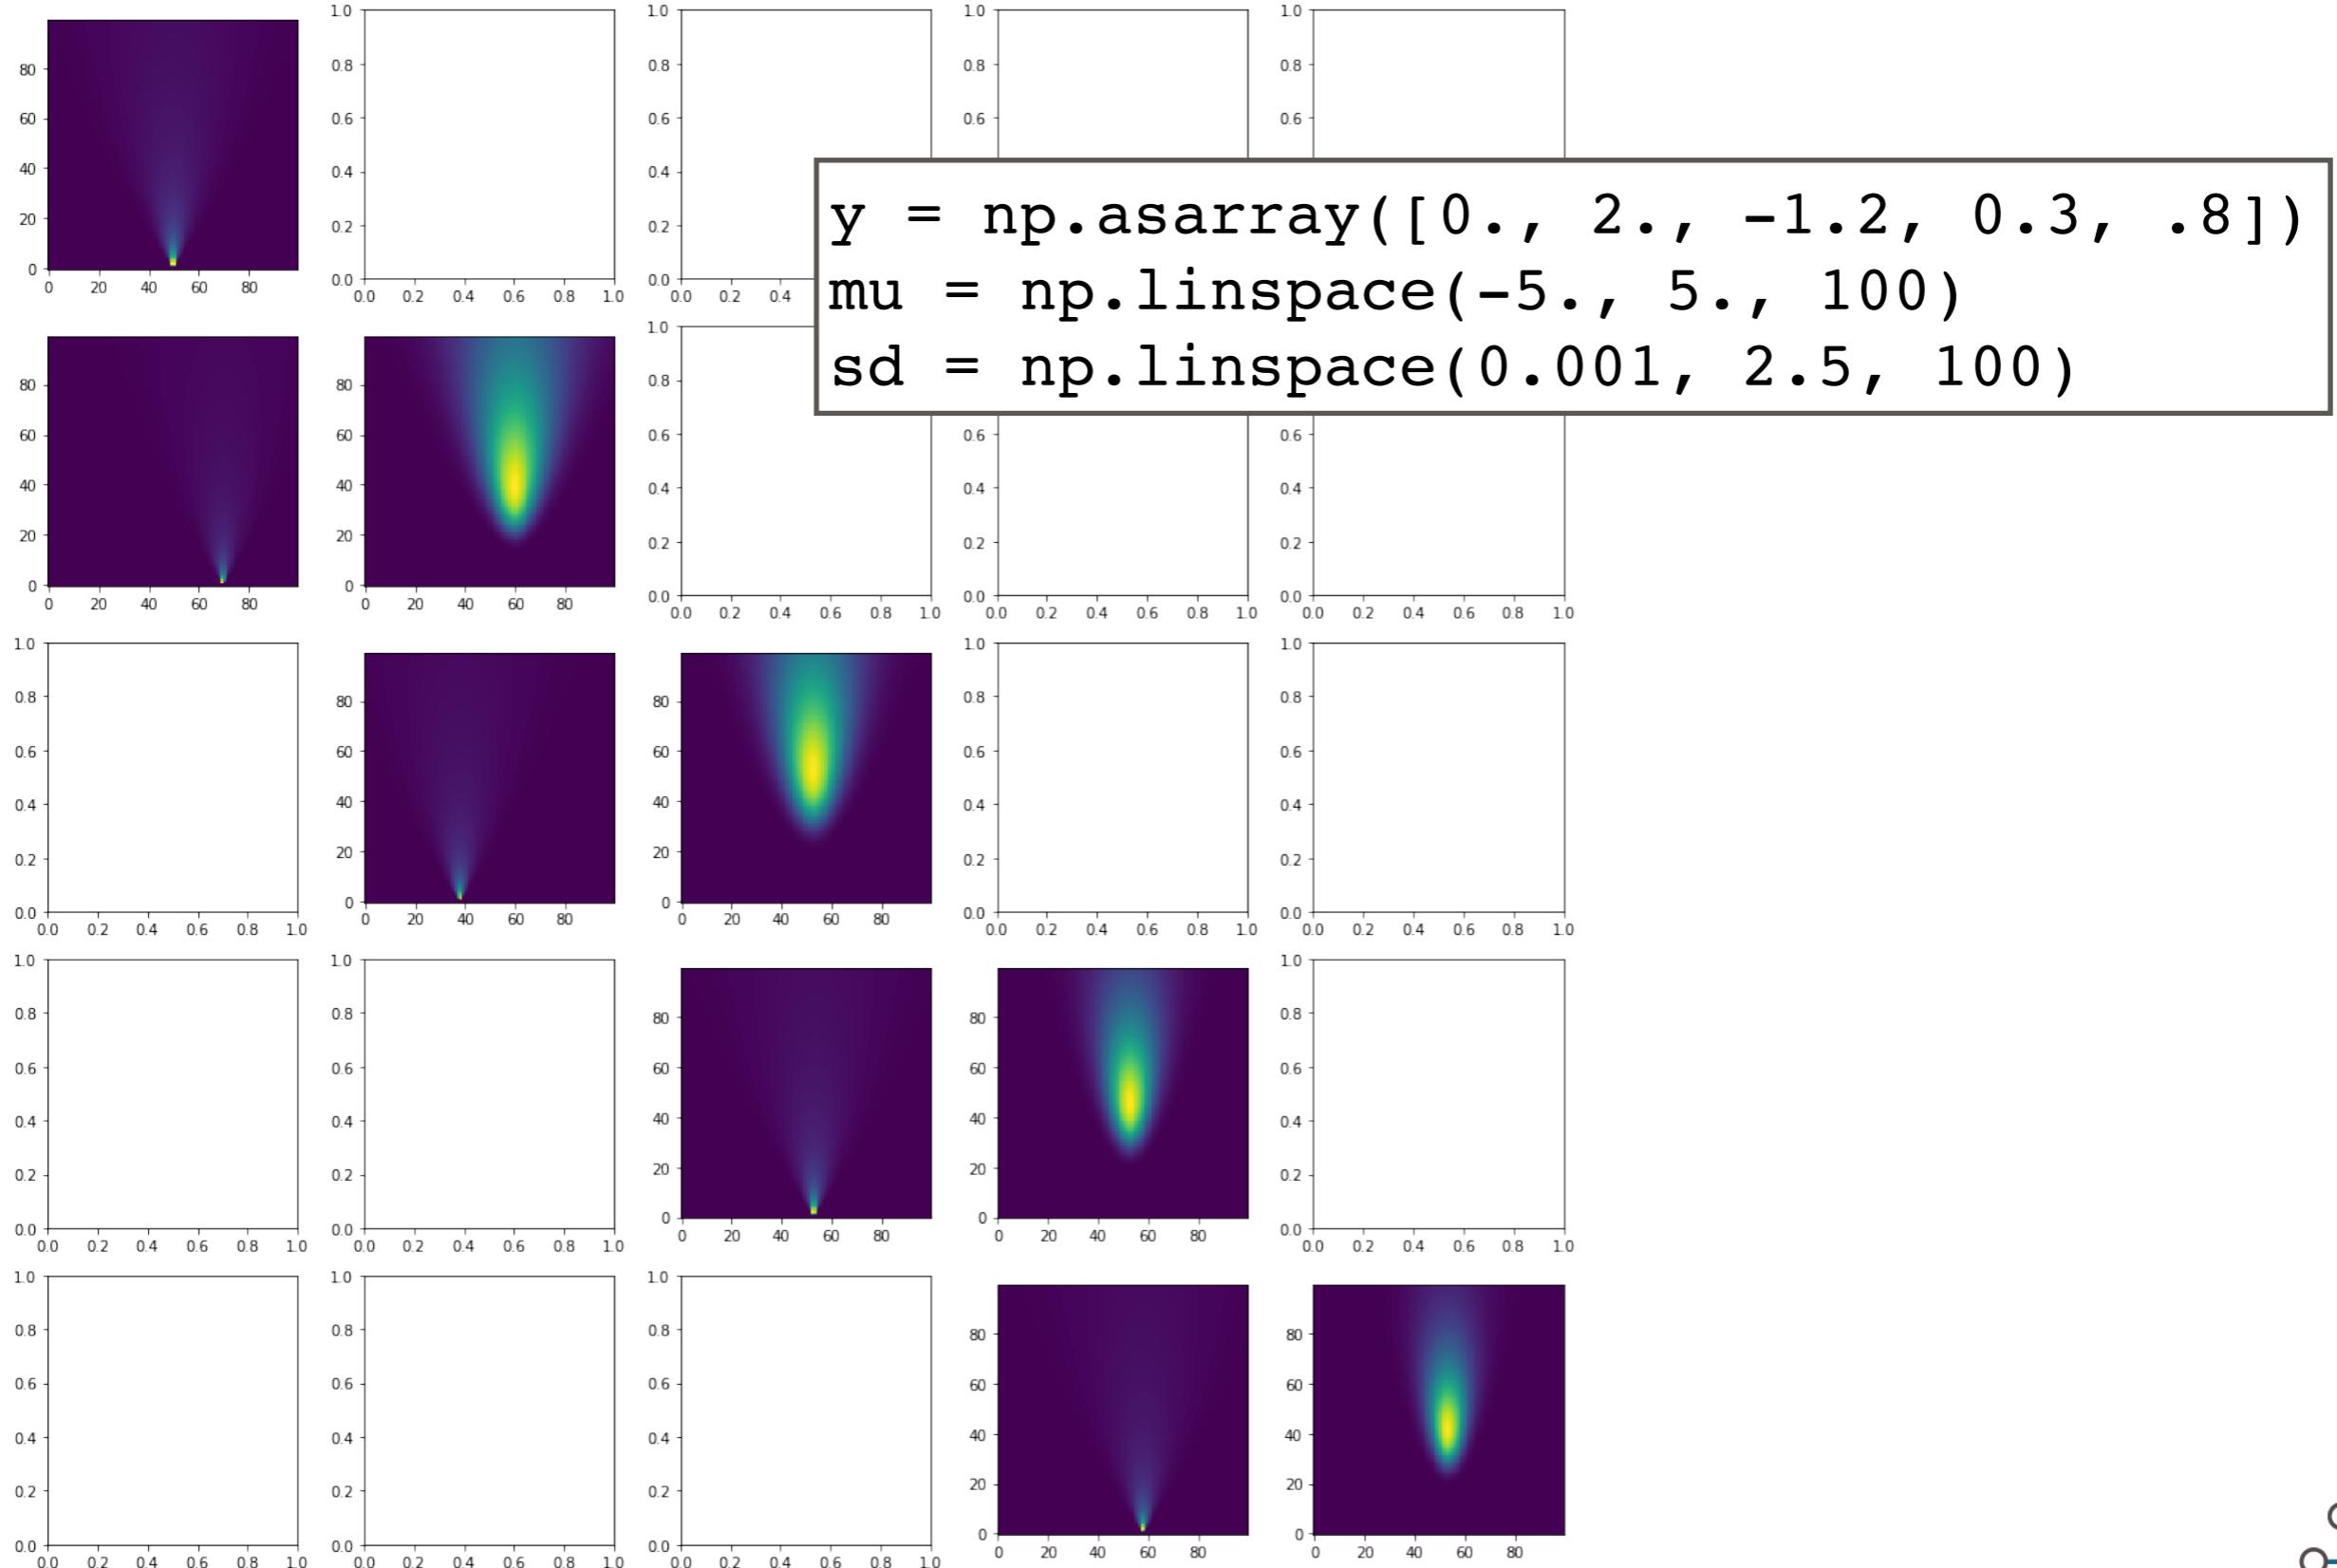
lambda p: st.binomial.pmf(y.sum(), len(y), p).prod()



Update, or cumulating information



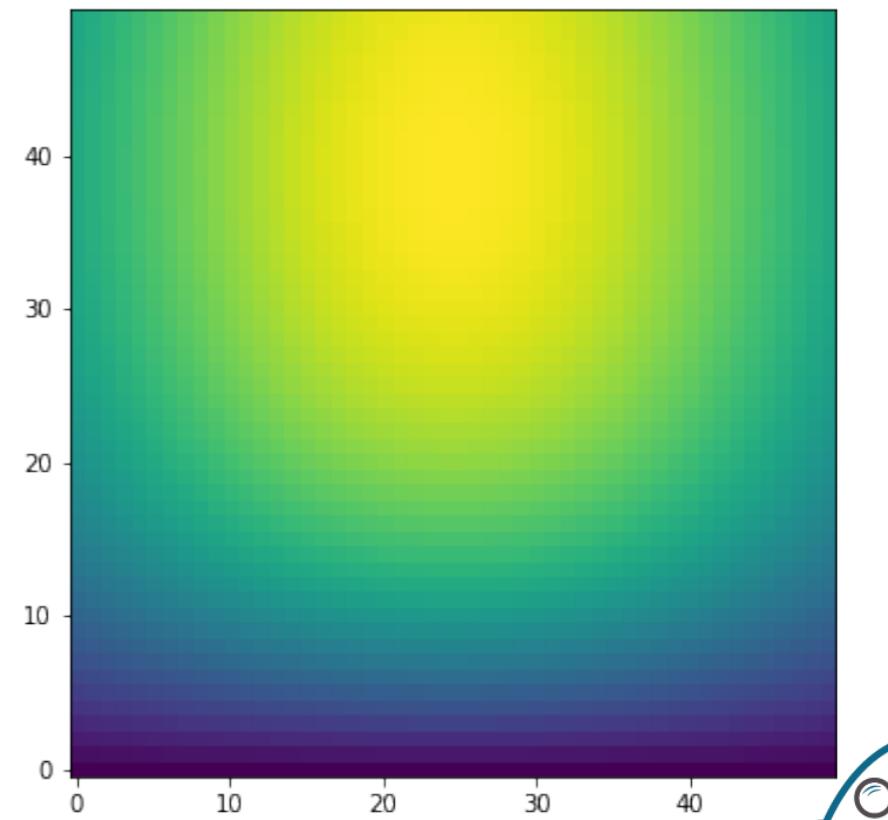
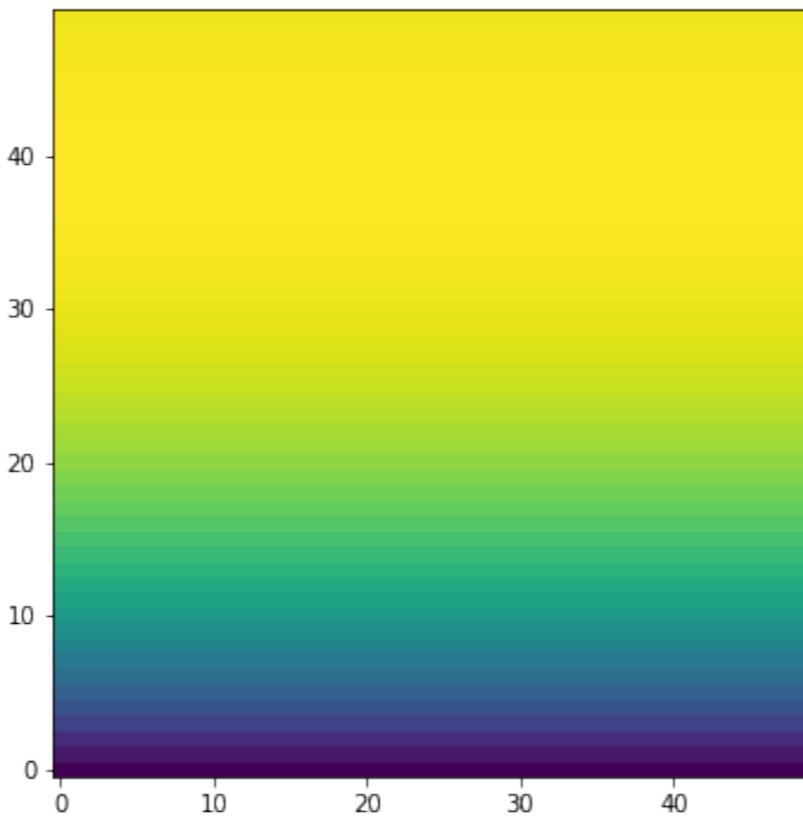
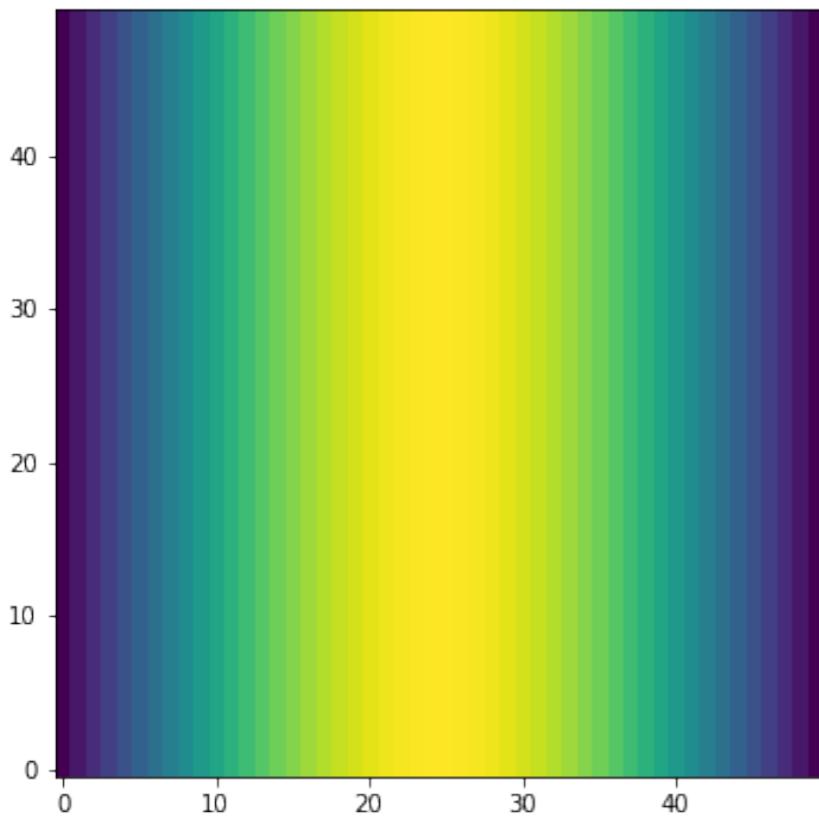
Update, or cumulating information



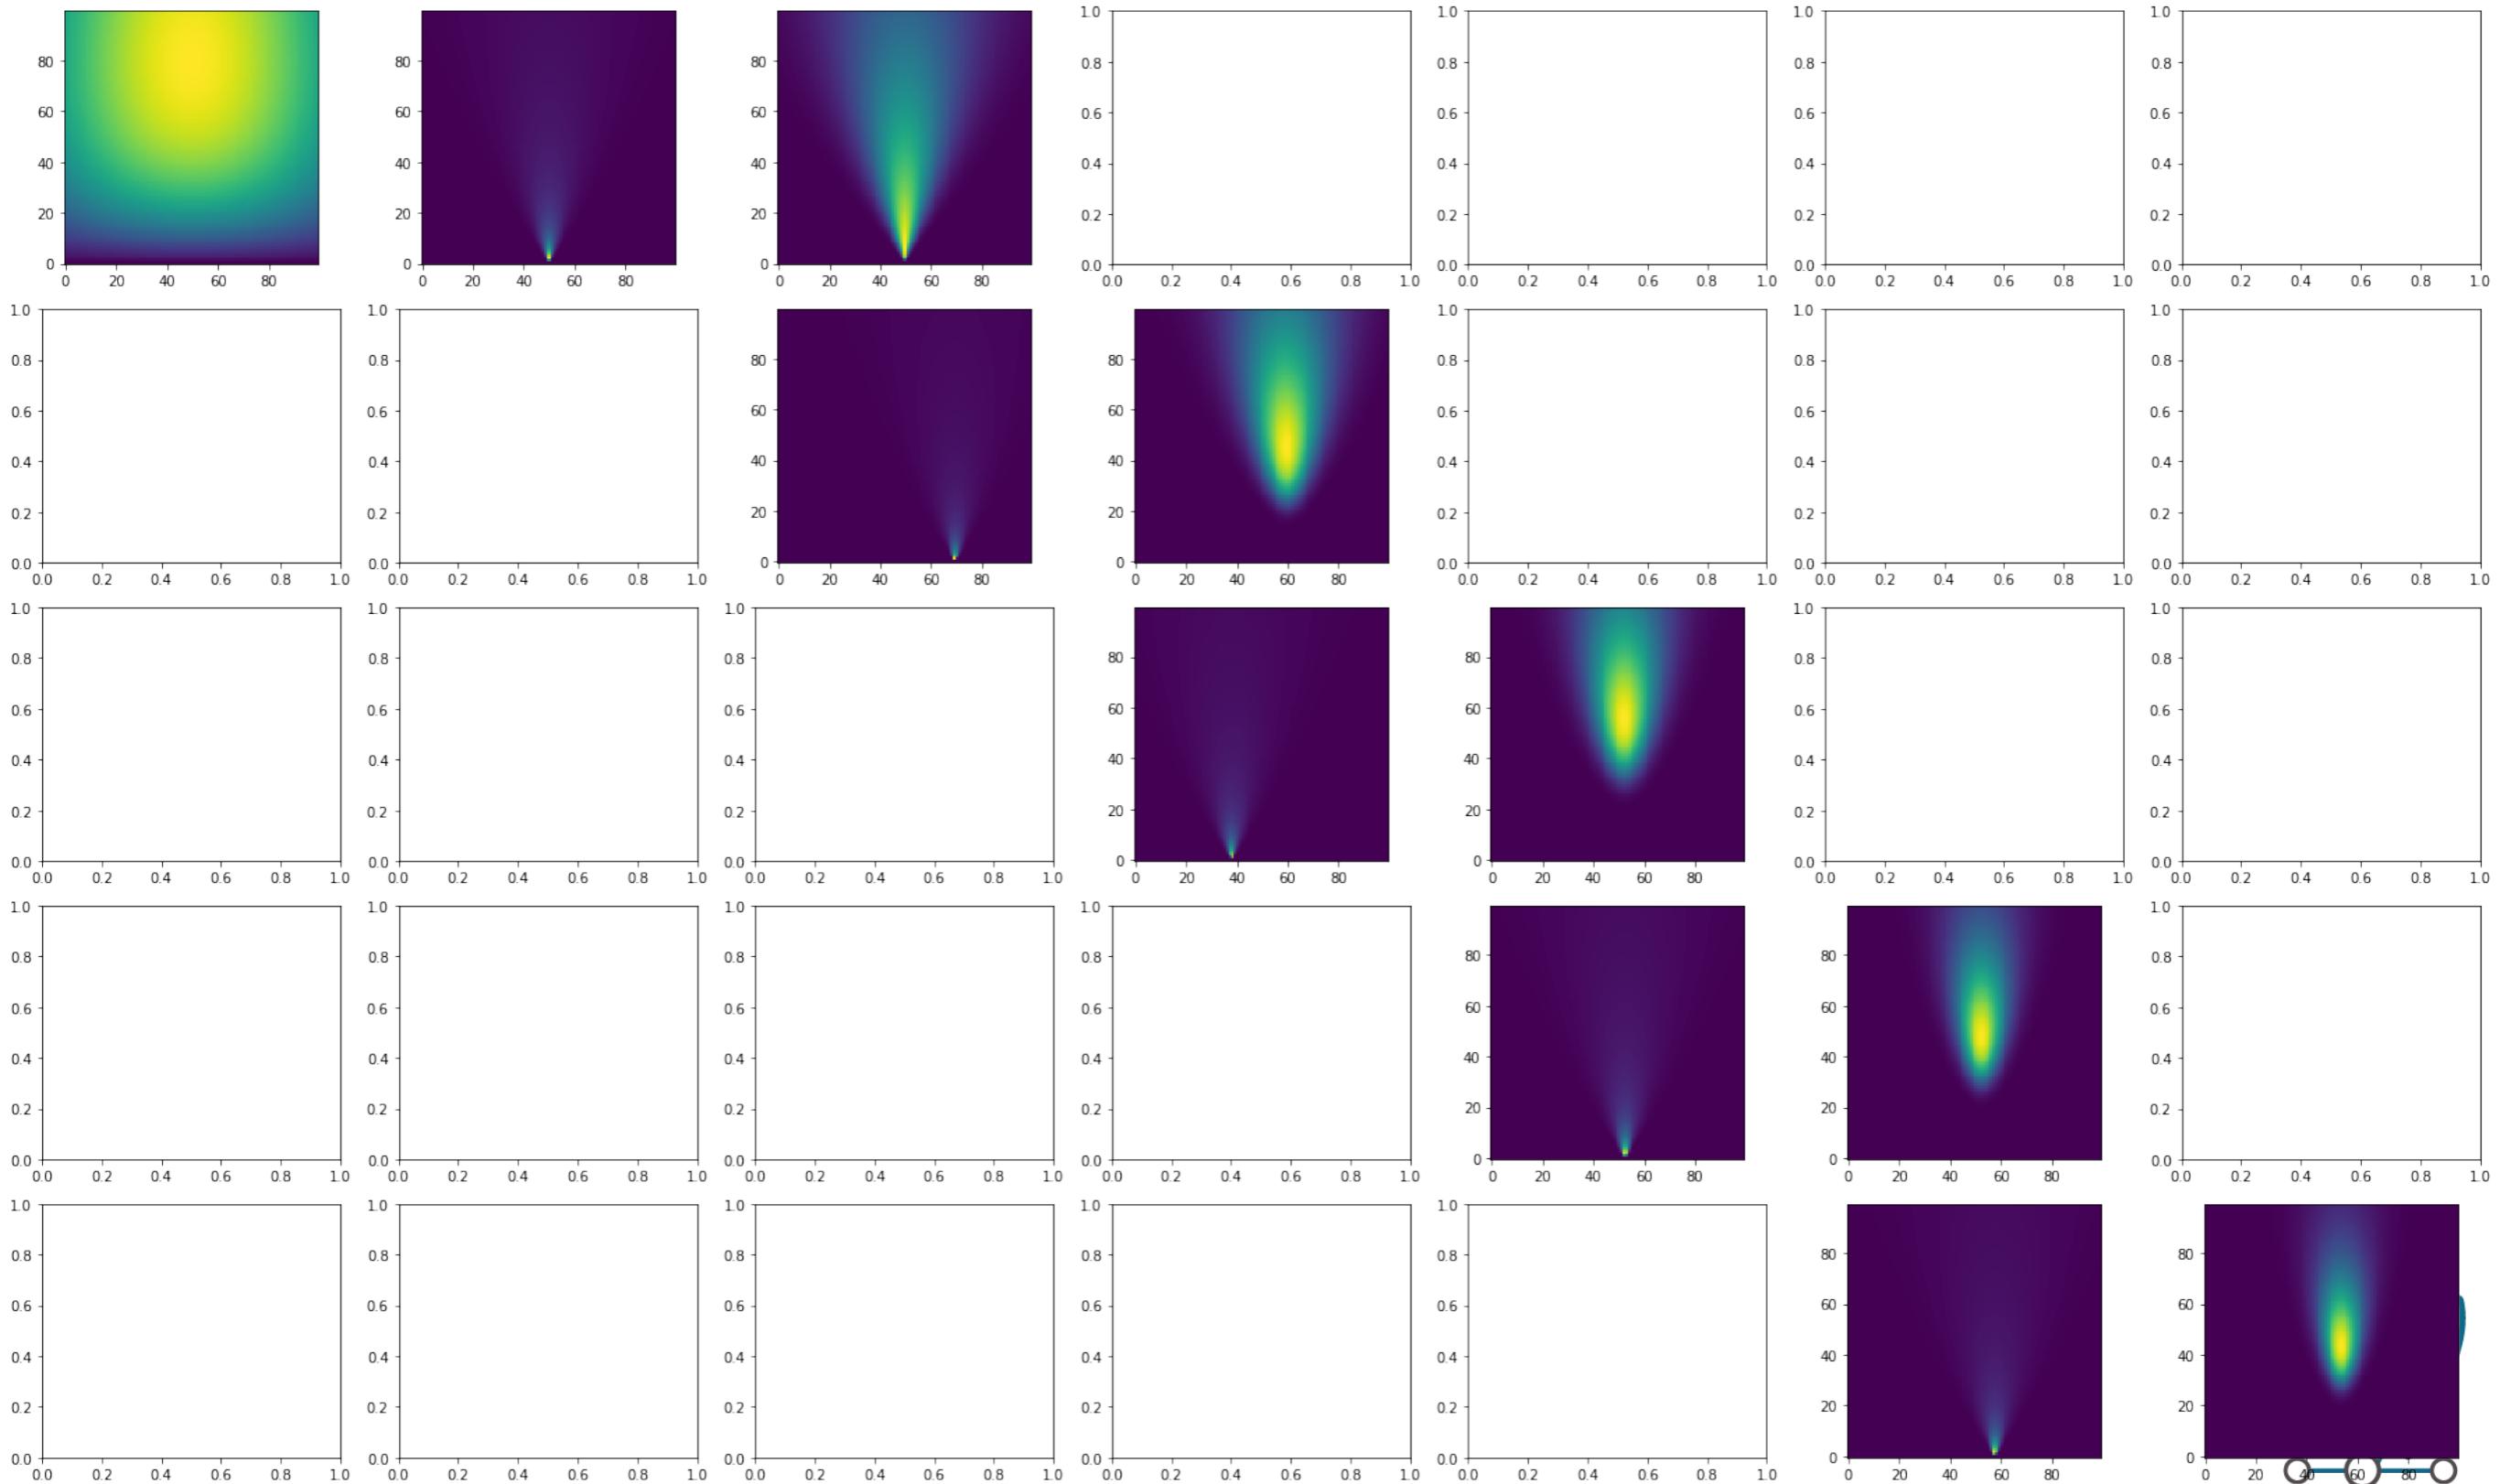
Priors

A and B	$P(A \cap B) = P(A B)P(B) = P(B A)P(A)$
	$P(A \cap B) = P(A)P(B)$ if A and B are independent

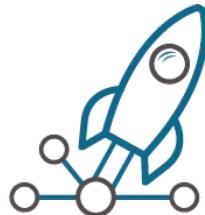
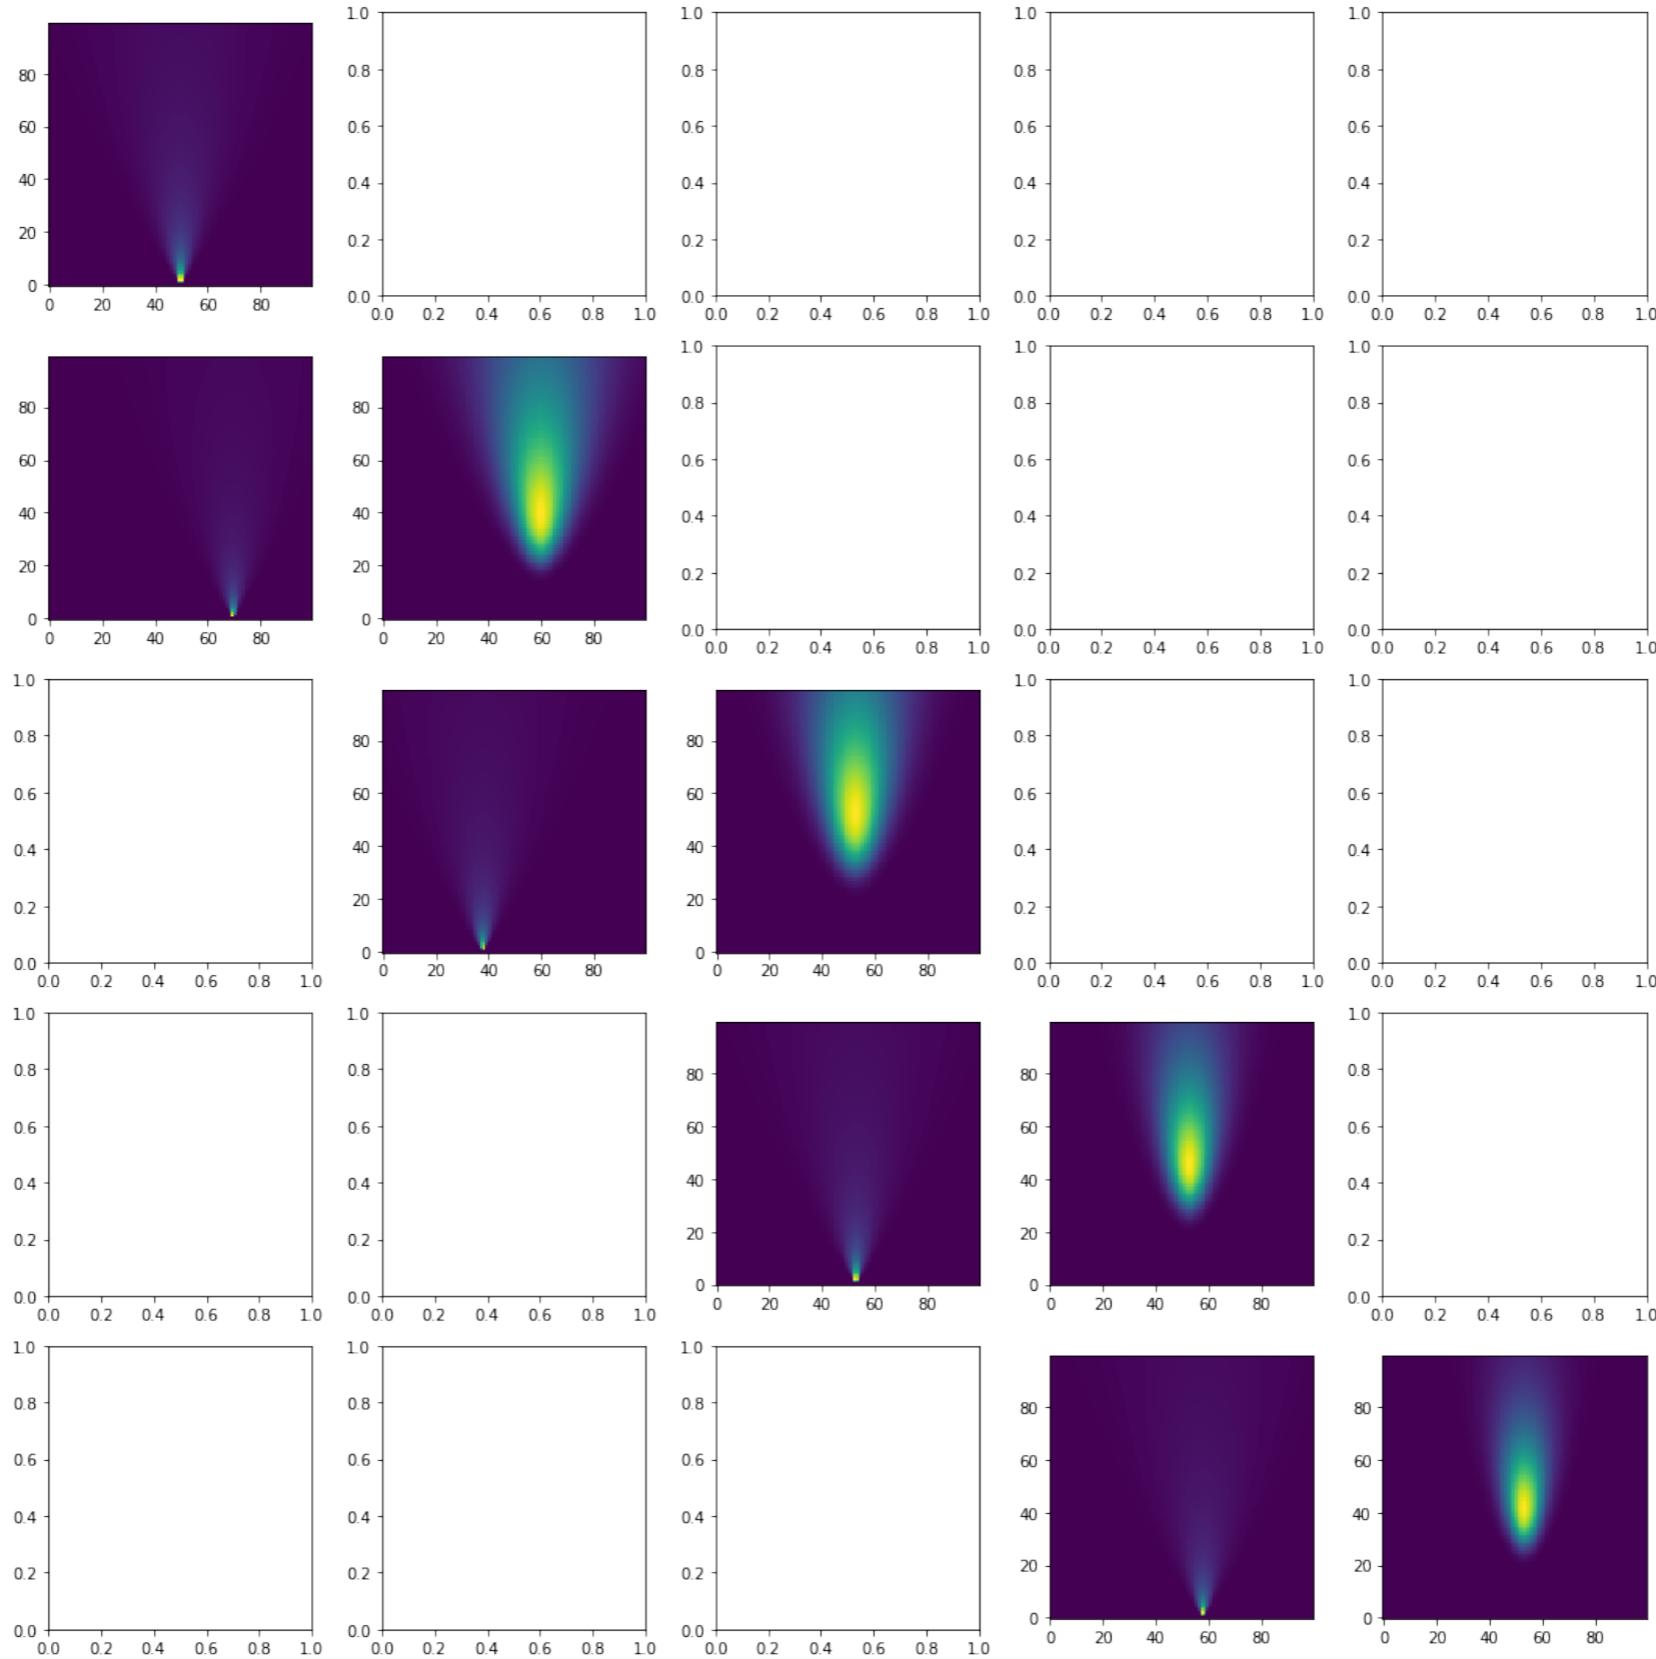
```
mu_prior = st.norm.pdf(mu, 0, 5)
sd_prior = st.gamma.pdf(sd, 2., scale=1.0/.5)
```



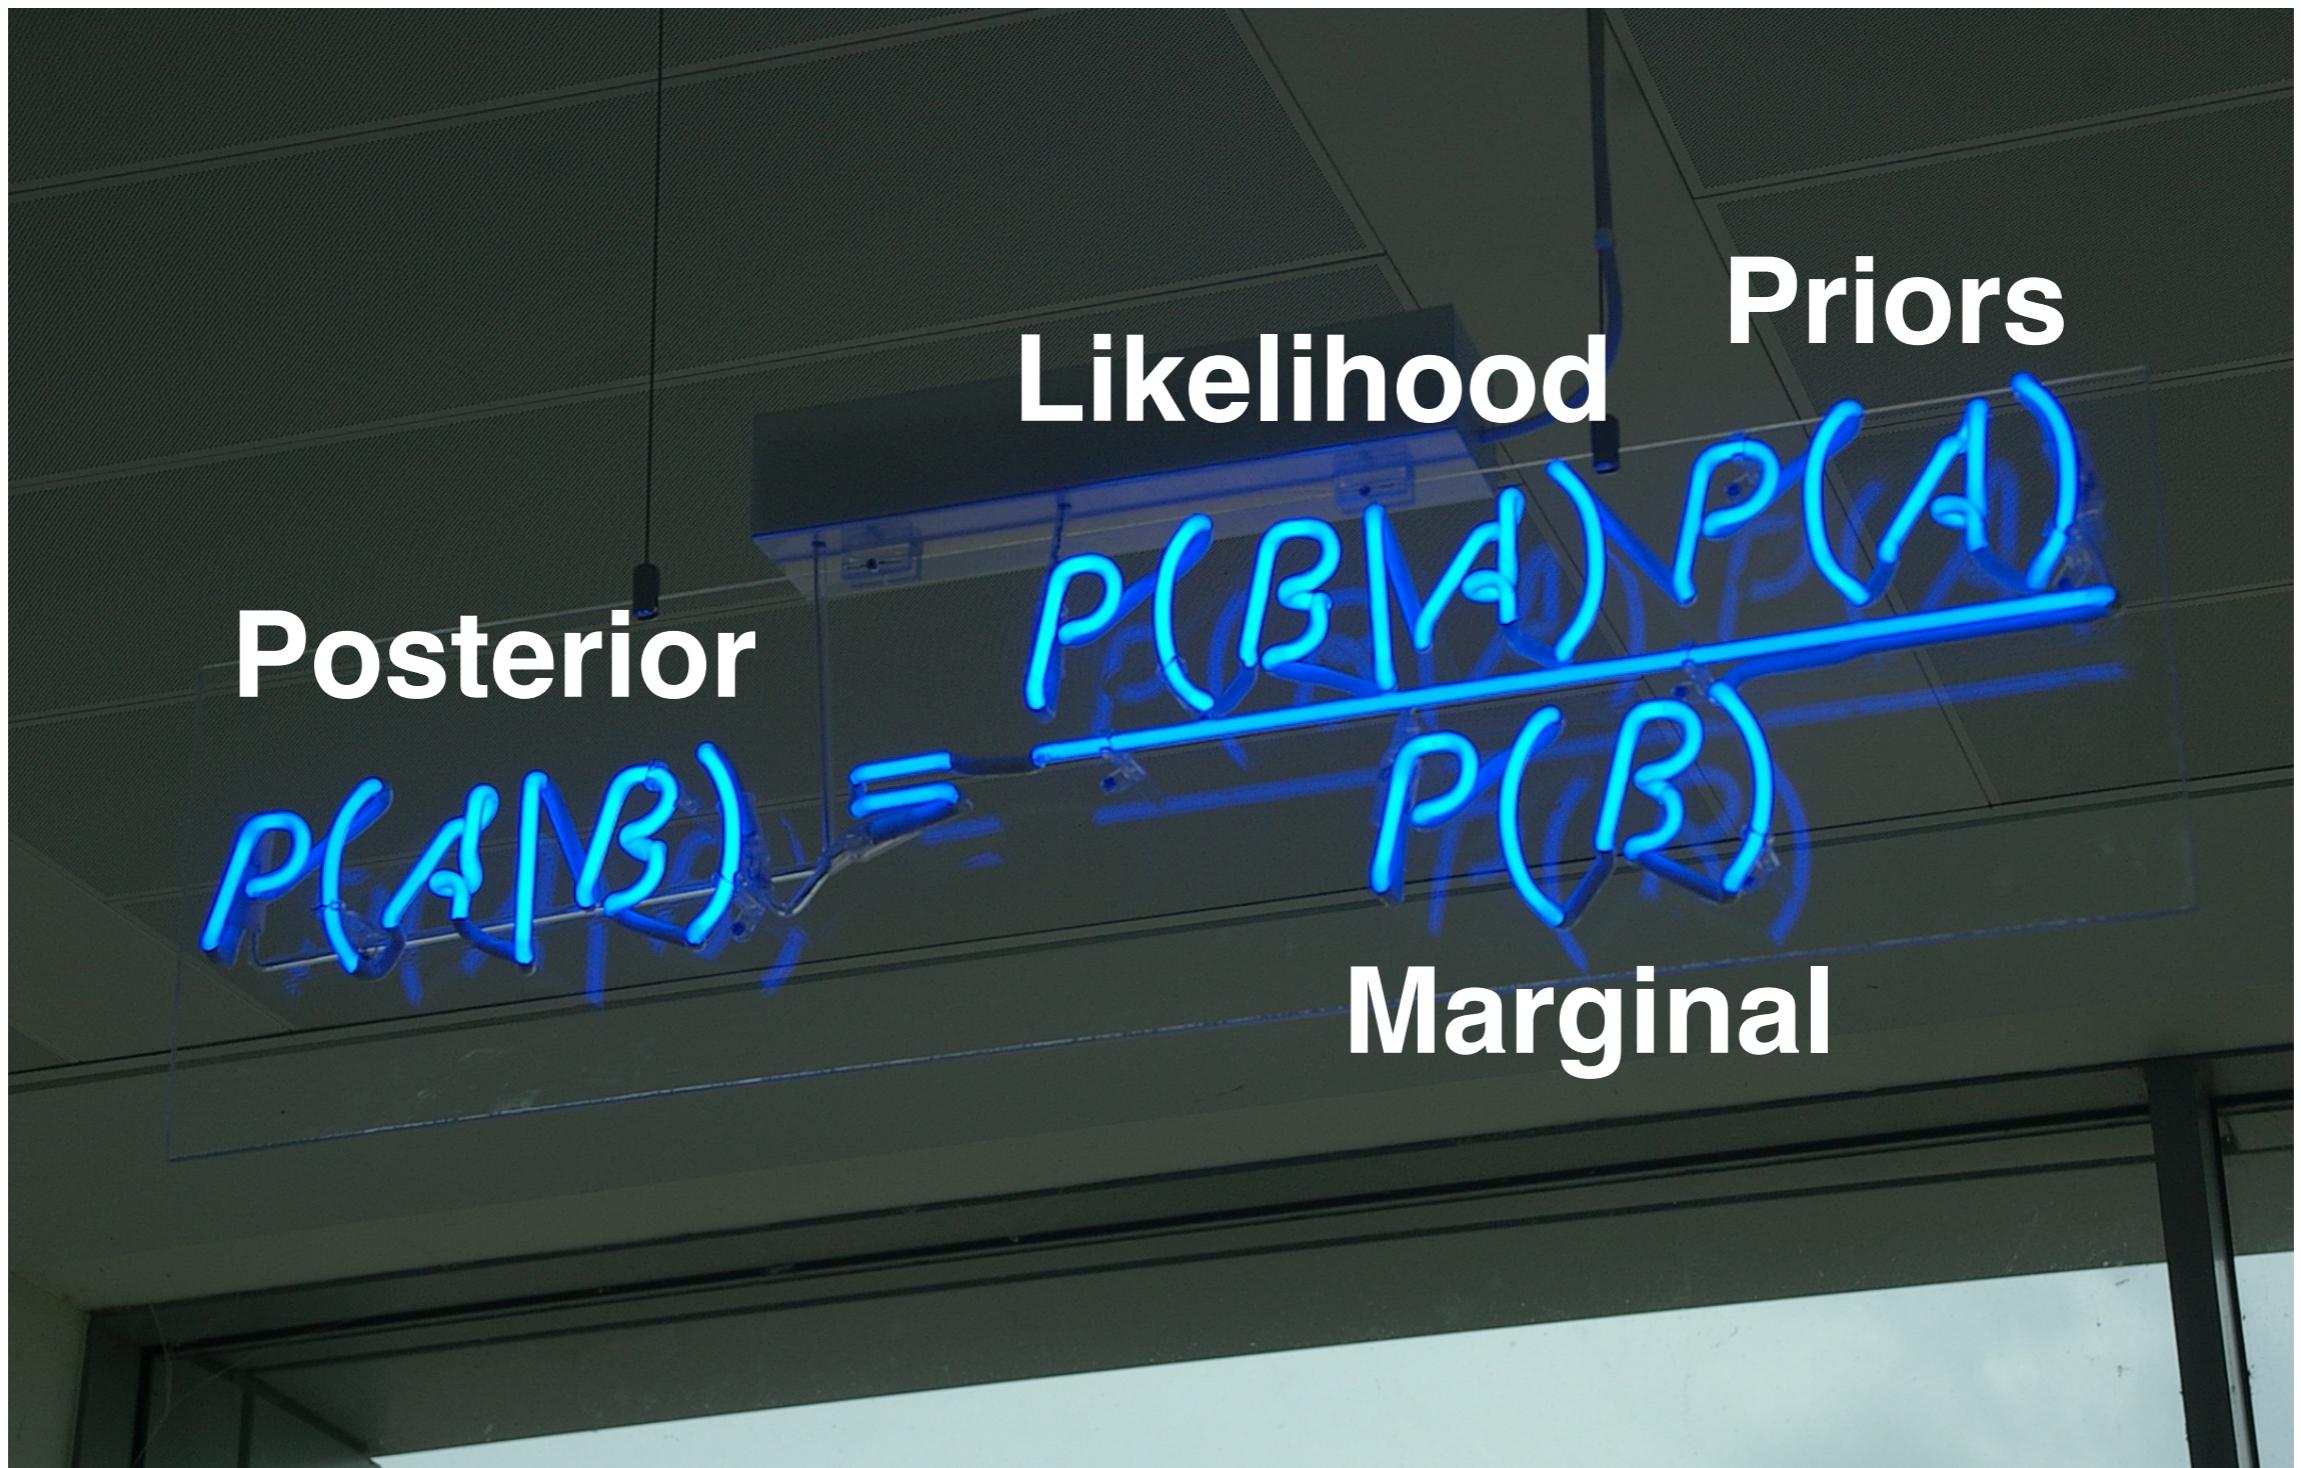
Update, or cumulating information



As comparison...



Bayes' theorem



Bayes's Rule

$$\begin{aligned} p(\theta|y) &= \frac{p(\theta, y)}{p(y)} && [\text{definition of conditional probability}] \\ &= \frac{p(y|\theta) p(\theta)}{p(y)} && [\text{chain rule}] \\ &= \frac{p(y|\theta) p(\theta)}{\int_{\Theta} p(y, \theta) d\theta} && [\text{law of total probability}] \\ &= \frac{p(y|\theta) p(\theta)}{\int_{\Theta} p(y|\theta) p(\theta) d\theta} && [\text{chain rule}] \\ &\propto p(y|\theta) p(\theta) && [y \text{ is fixed}] \end{aligned}$$



Bayesian models

$$p(\text{Parameters} \mid \text{Data}) \propto p(\text{Data} \mid \text{Parameters}) * p(\text{Parameters})$$

$$\mu \sim \text{Normal}(0, 5)$$

$$\sigma \sim \text{Gamma}(2, 5)$$

$$y \sim \text{Normal}(\mu, \sigma)$$

$$\alpha \sim \text{Gamma}(1, 1)$$

$$\beta_1, \dots, \beta_K \sim \text{Beta}(1, \alpha)$$

$$w_i = \beta_i \prod_{j=i-1}^i (1 - \beta_j)$$

$$\lambda_1, \dots, \lambda_K \sim U(0, 5)$$

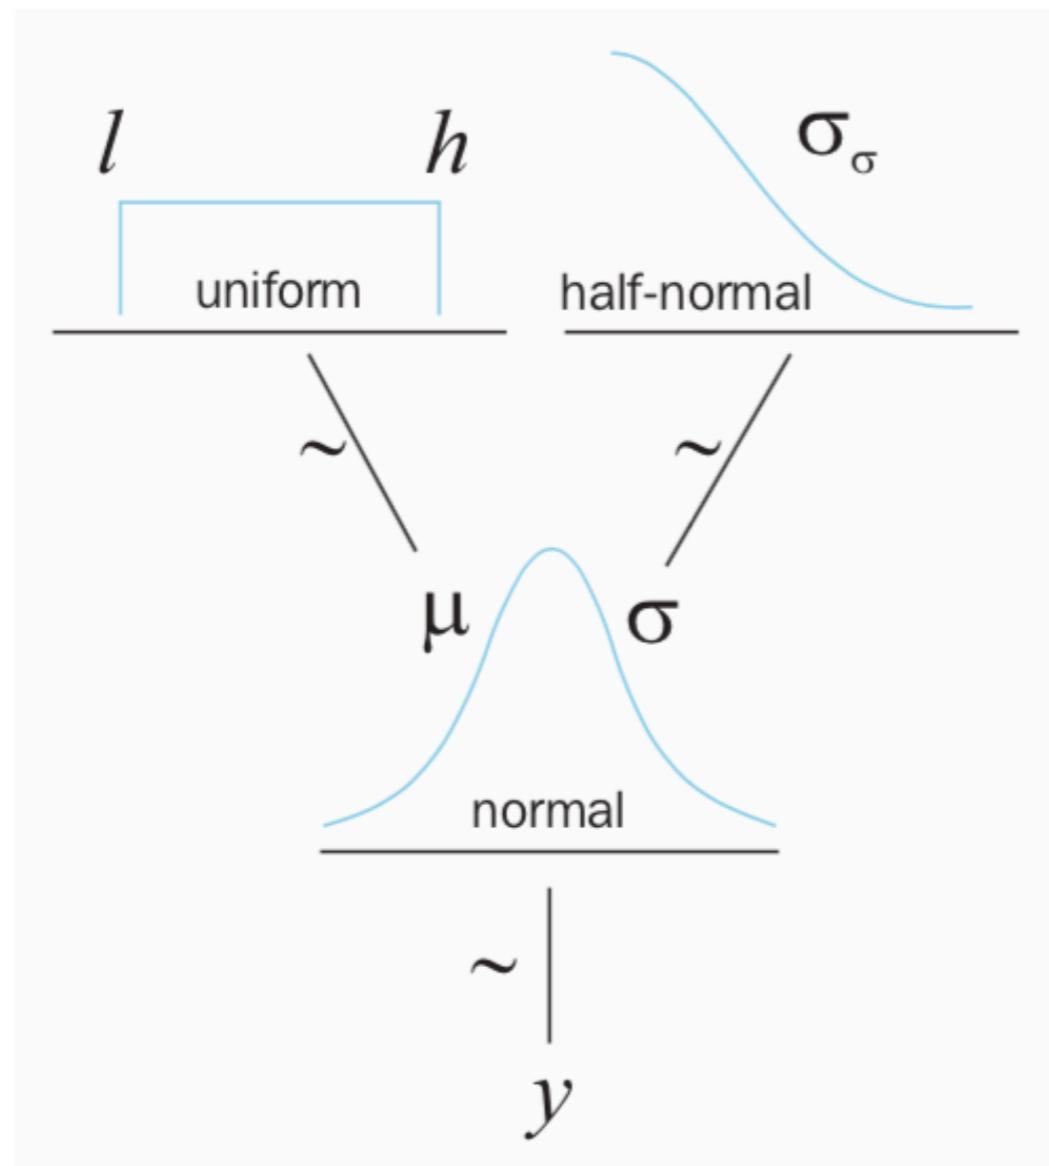
$$\tau_1, \dots, \tau_K \sim \text{Gamma}(1, 1)$$

$$\mu_i \mid \lambda_i, \tau_i \sim N(0, (\lambda_i \tau_i)^{-1})$$

$$x \mid w_i, \lambda_i, \tau_i, \mu_i \sim \sum_{i=1}^K w_i N(\mu_i, (\lambda_i \tau_i)^{-1})$$



Generative model

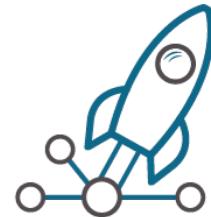
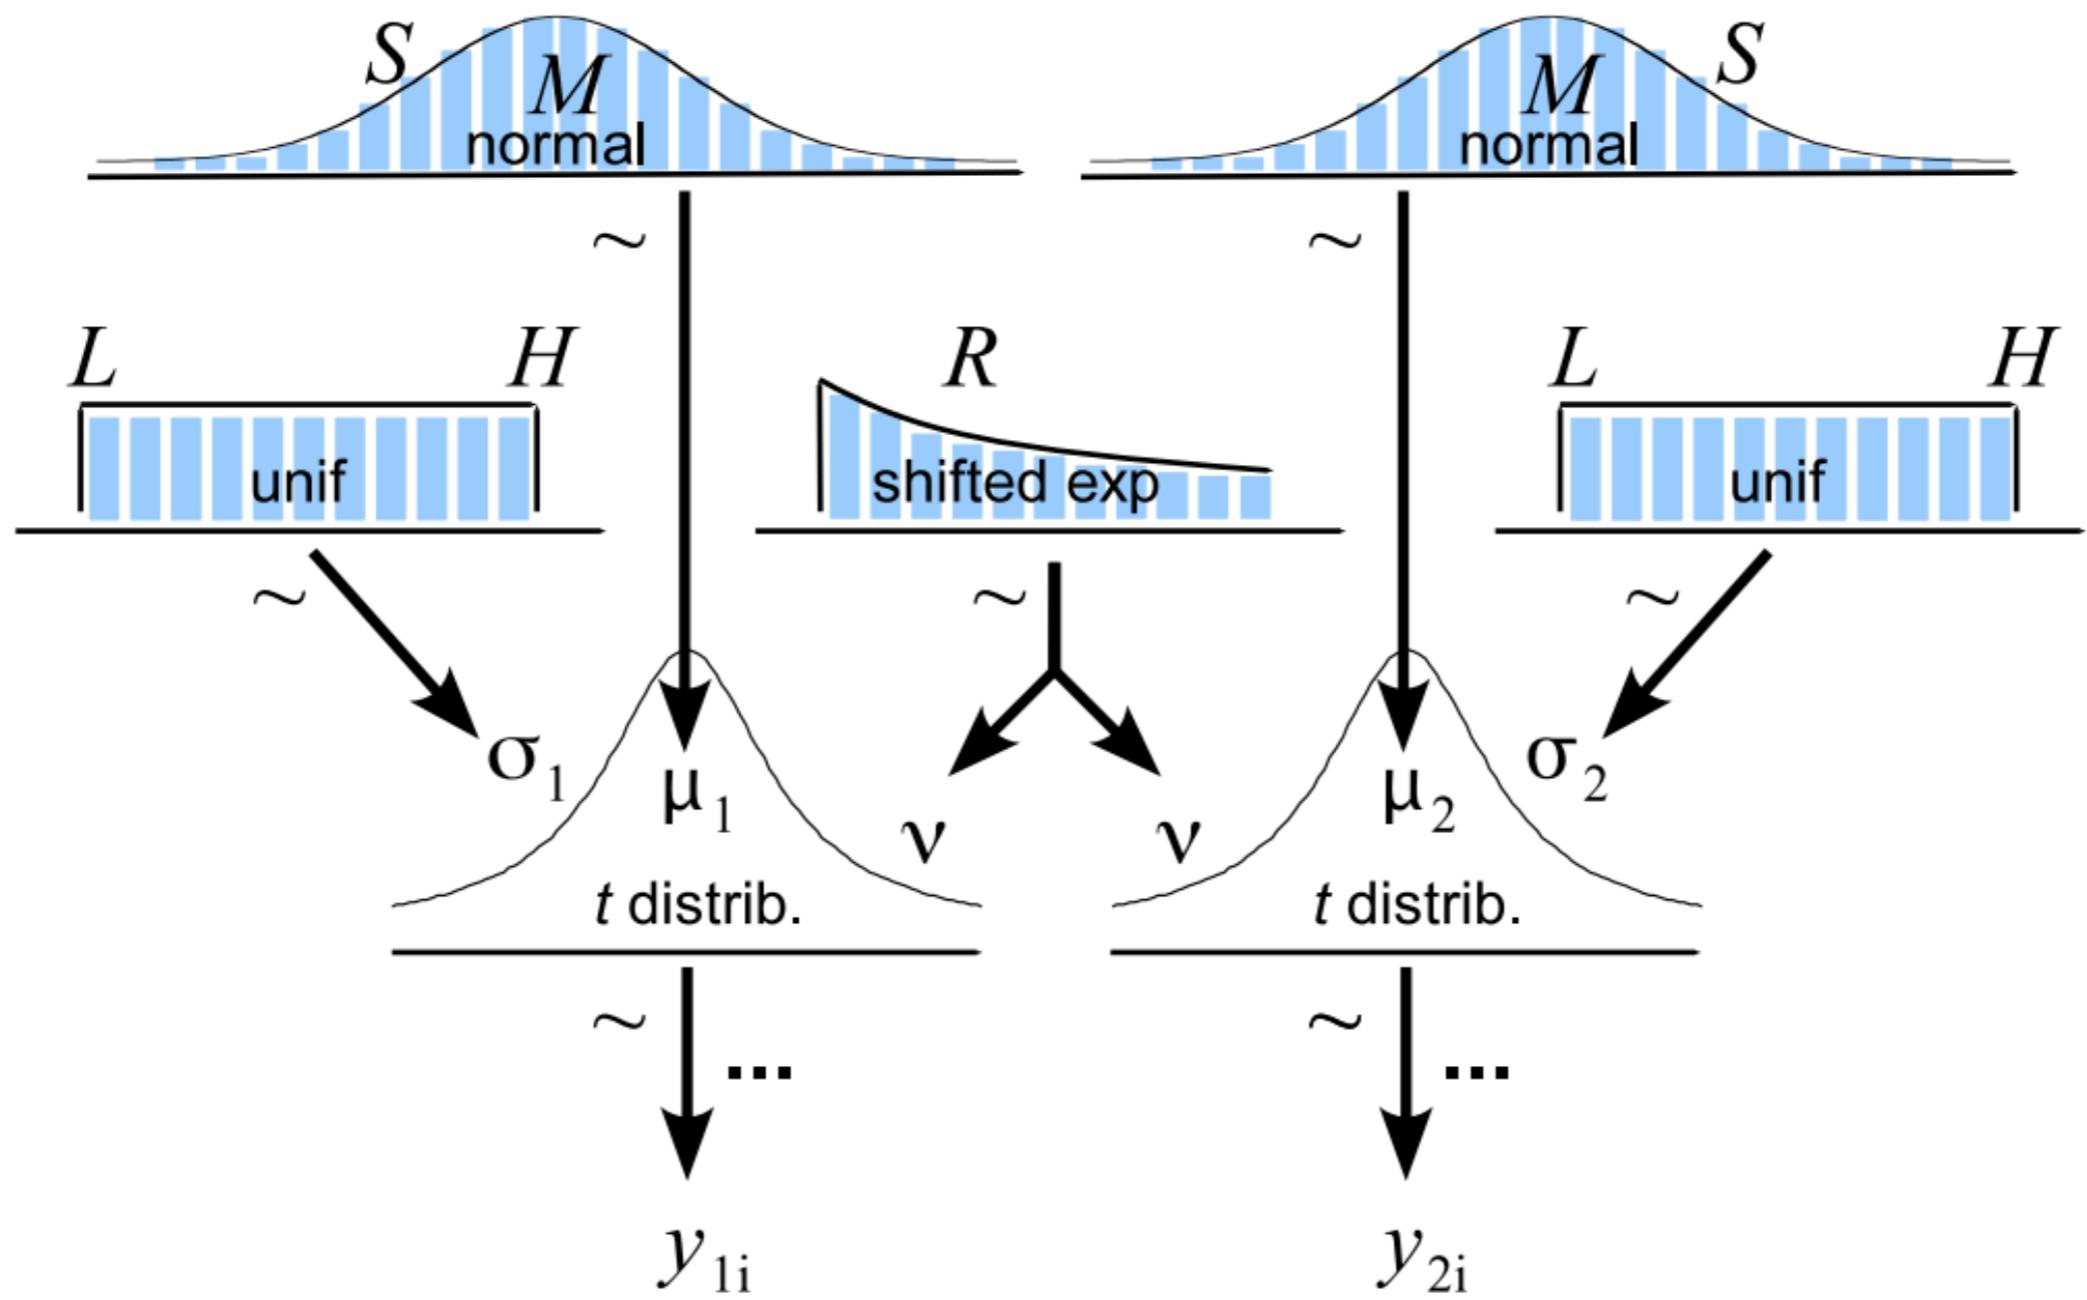


$\text{mu} \sim \text{Uniform}(-5, 5)$
 $\text{sd} \sim \text{HalfNormal}(2.5)$
 $y \sim \text{Normal}(\text{mu}, \text{sd})$

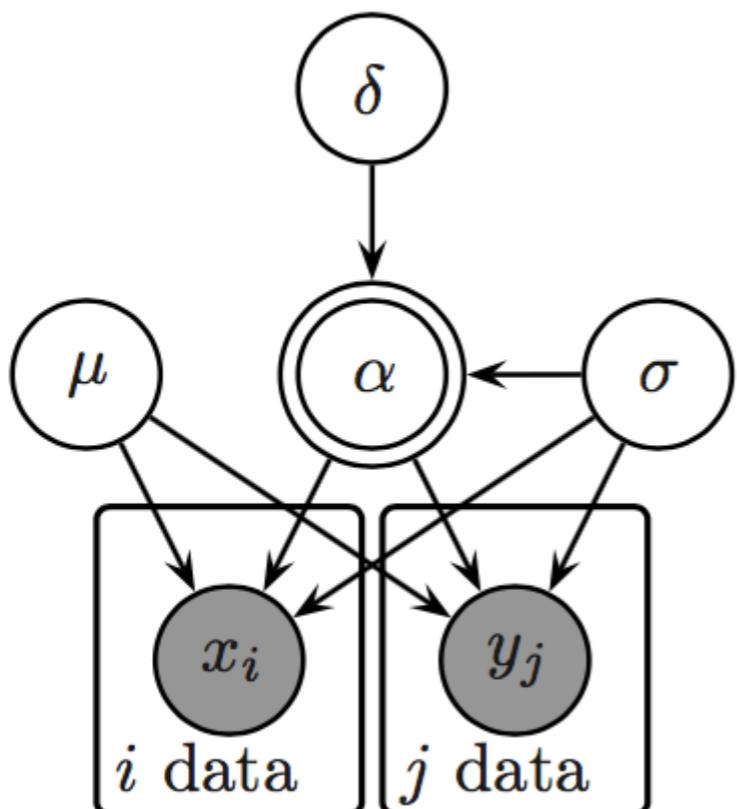


BEST

KRUSCHKE



Bayesian Graphical Models



$$\delta \sim \text{Cauchy}(0, 1)$$

$$\mu \sim \text{Cauchy}(0, 1)$$

$$\sigma \sim \text{Cauchy}(0, 1)_{\mathcal{I}(0, \infty)}$$

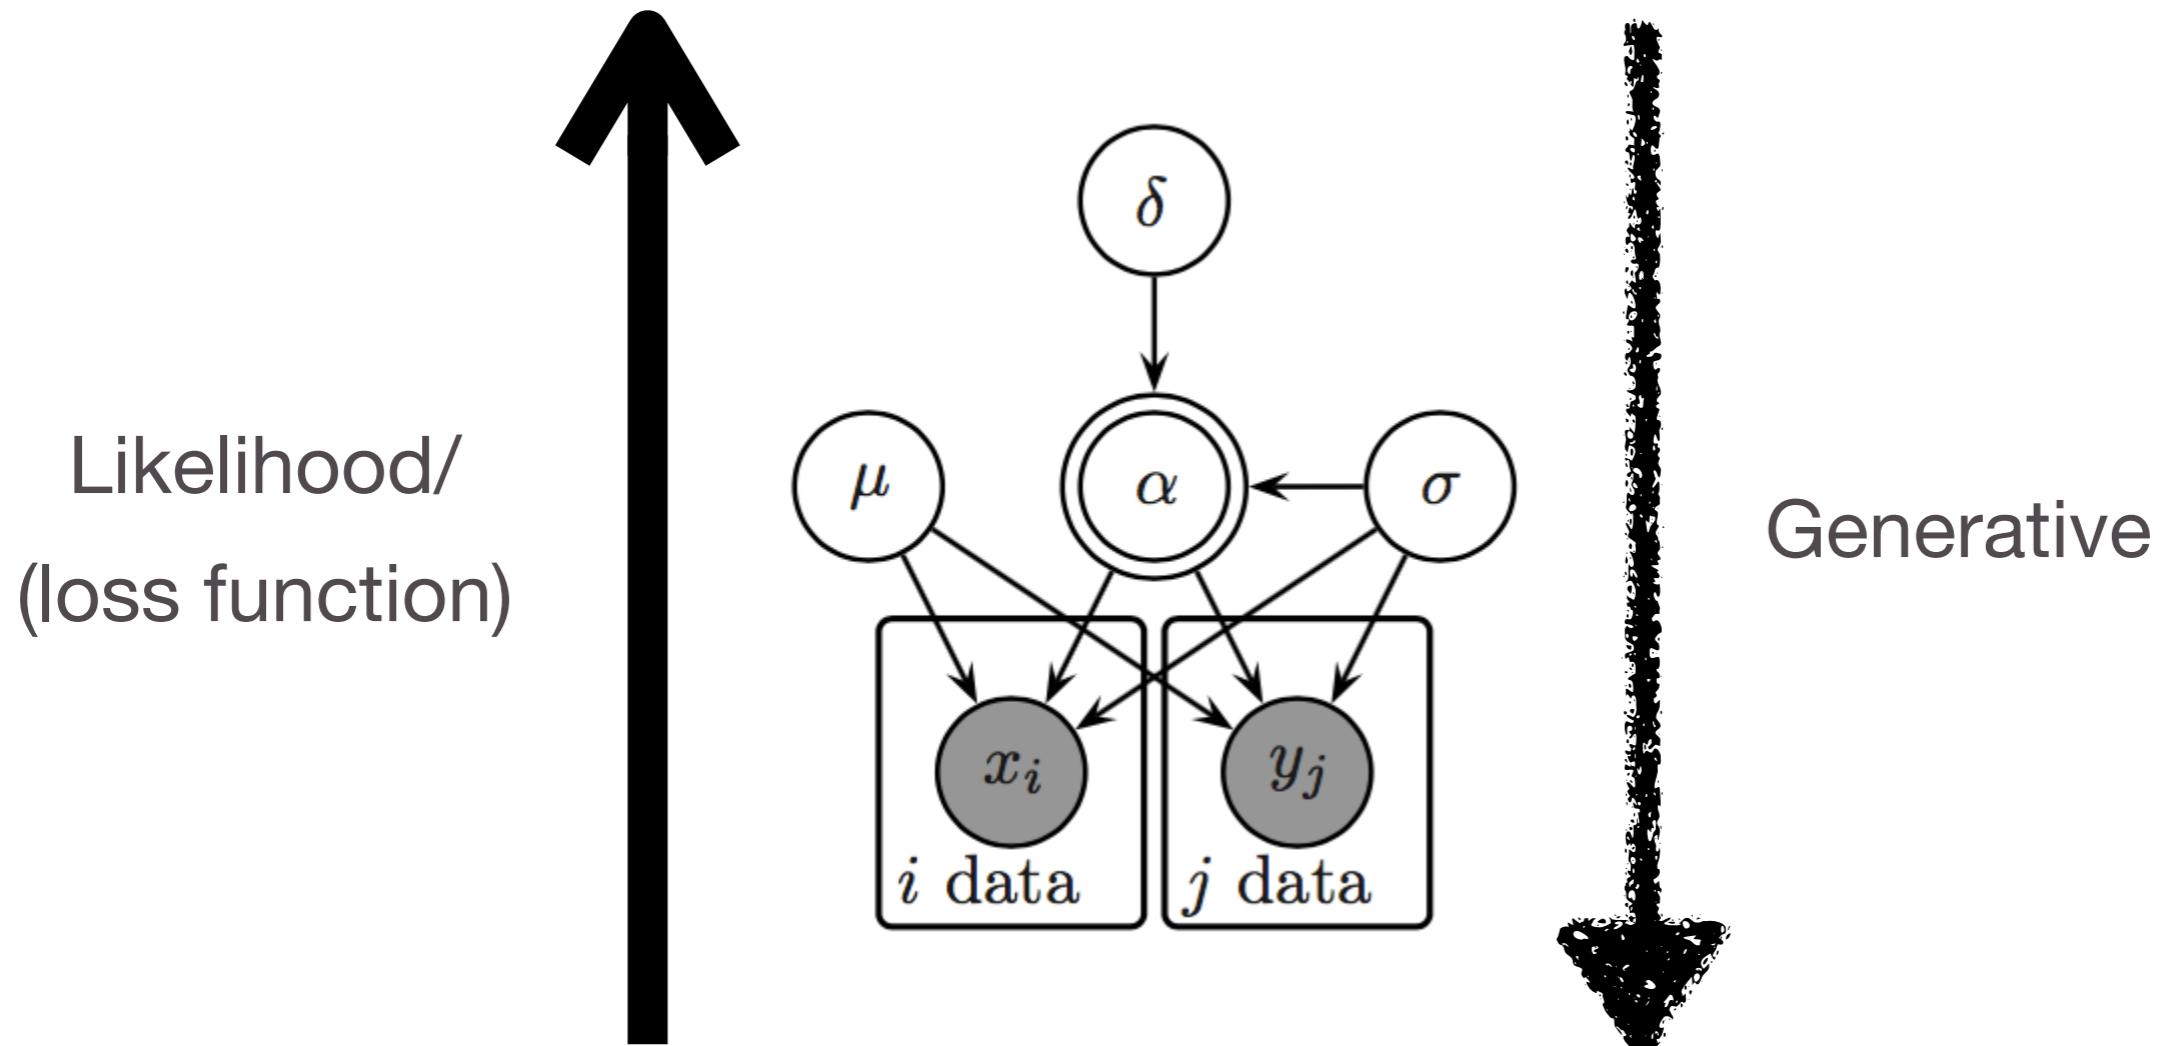
$$\alpha \leftarrow \delta\sigma$$

$$x_i \sim \text{Gaussian}(\mu + \frac{\alpha}{2}, 1/\sigma^2)$$

$$y_j \sim \text{Gaussian}(\mu - \frac{\alpha}{2}, 1/\sigma^2)$$



Bayesian Graphical Models



Components in PPL

```
...
739     def __init__(self, lam, *args, **kwargs):
740         super(Exponential, self).__init__(*args, **kwargs)
741         self.lam = lam = tt.as_tensor_variable(lam)
742         self.mean = 1. / self.lam
743         self.median = self.mean * tt.log(2)
744         self.mode = tt.zeros_like(self.lam)
745
746         self.variance = self.lam**-2
747
748         assert_negative_support(lam, 'lam', 'Exponential')
749
750     def random(self, point=None, size=None):
751         lam = draw_values([self.lam], point=point)[0]
752         return generate_samples(np.random.exponential, scale=1. / lam,
753                               dist_shape=self.shape,
754                               size=size)
755
756     def logp(self, value):
757         lam = self.lam
758         return bound(tt.log(lam) - lam * value, value >= 0, lam > 0)
759
```



Components in PPL

...

```
47 _DISTRIBUTION_PUBLIC_METHOD_WRAPPERS = [
48     "batch_shape",
49     "batch_shape_tensor",
50     "cdf",
51     "covariance",
52     "cross_entropy",
53     "entropy",
54     "event_shape",
55     "event_shape_tensor",
56     "kl_divergence",
57     "log_cdf",
58     "log_prob",
59     "log_survival_function",
60     "mean",
61     "mode",
62     "prob",
63     "sample",
64     "stddev",
65     "survival_function",
66     "variance",
67 ]
68
```

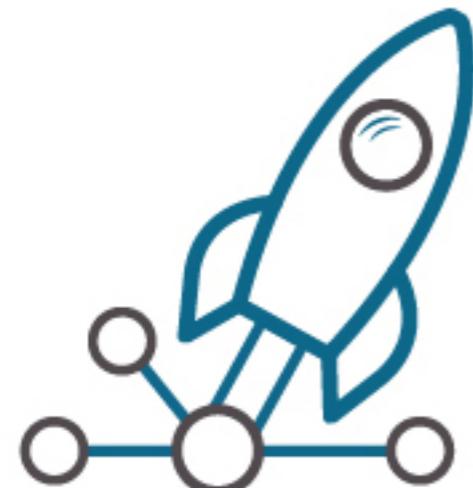
TensorFlow distributions





From the [PyMC3 documentation](#):

PyMC3 is a Python package for Bayesian statistical modeling and Probabilistic Machine Learning which focuses on **advanced Markov chain Monte Carlo** and variational fitting algorithms. Its flexibility and extensibility make it applicable to a large suite of problems.



PyMC



TensorFlow

From the [PyMC3 documentation](#):

PyMC3 is a Python package for Bayesian statistical modeling and Probabilistic Machine Learning which focuses on **advanced Markov chain Monte Carlo** and variational fitting algorithms. Its flexibility and extensibility make it applicable to a large suite of problems.

Why PyMC3?

Easy to interact with different components of your model.



Coin flip in PyMC3

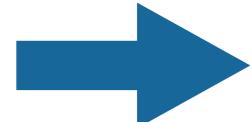
$k = 7$

$k \sim \text{Binomial}(n=n, p=p)$

$n \sim \text{DiscreteUniform}(0, 25)$

$p \sim \text{Uniform}(0., 1.)$

```
with pm.Model() as m:  
    n = pm.DiscreteUniform('n', 0, 25)  
    p = pm.Uniform('p', 0., 1., transform=None)  
    y = pm.Binomial('k', n, p, observed=7)
```



Code3 - Combining_Likelihood.ipynb



Coin flip in PyMC3

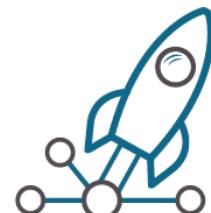
```
with pm.Model() as m:  
    n = pm.DiscreteUniform('n', 0, 25)  
    p = pm.Uniform('p', 0., 1., transform=None)  
    y = pm.Binomial('k', n, p, observed=7)
```

```
point = m.test_point  
{'n': array(12), 'p': array(0.5)}
```

```
logp_m = m.logp  
logp_y = y.logp
```

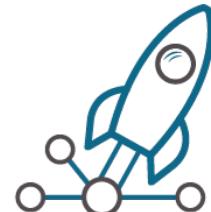
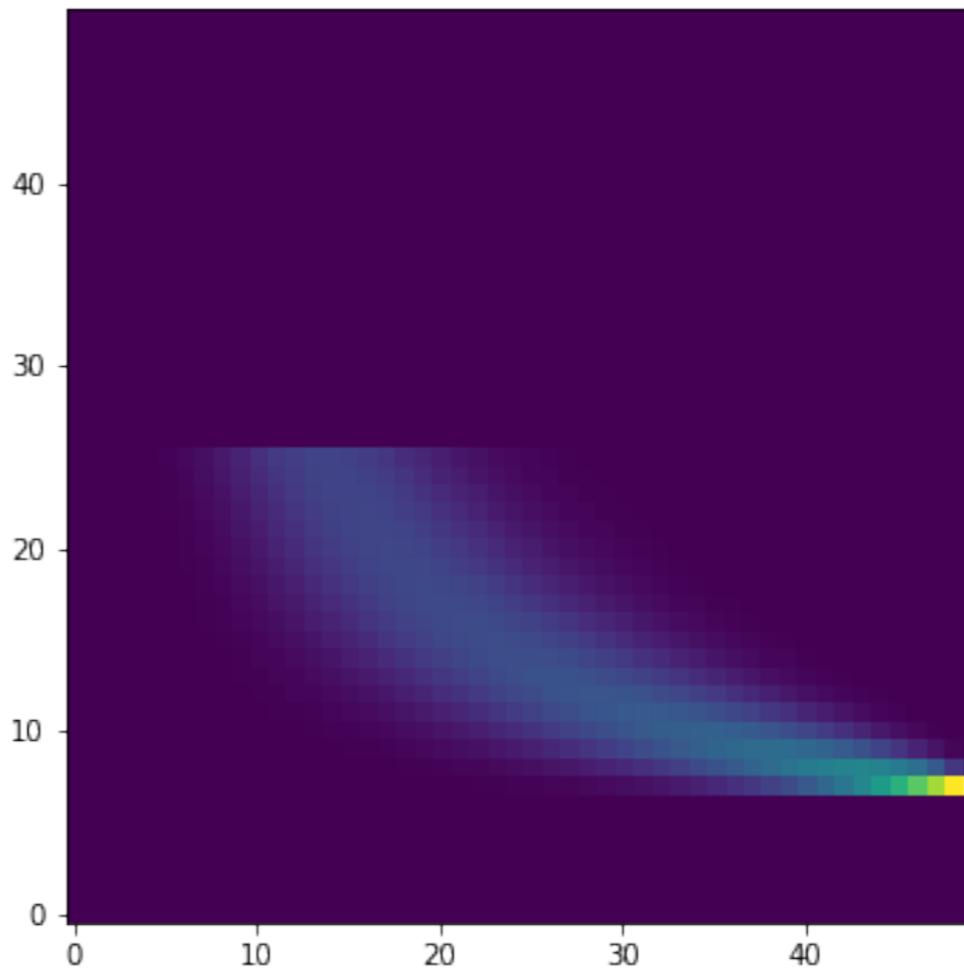
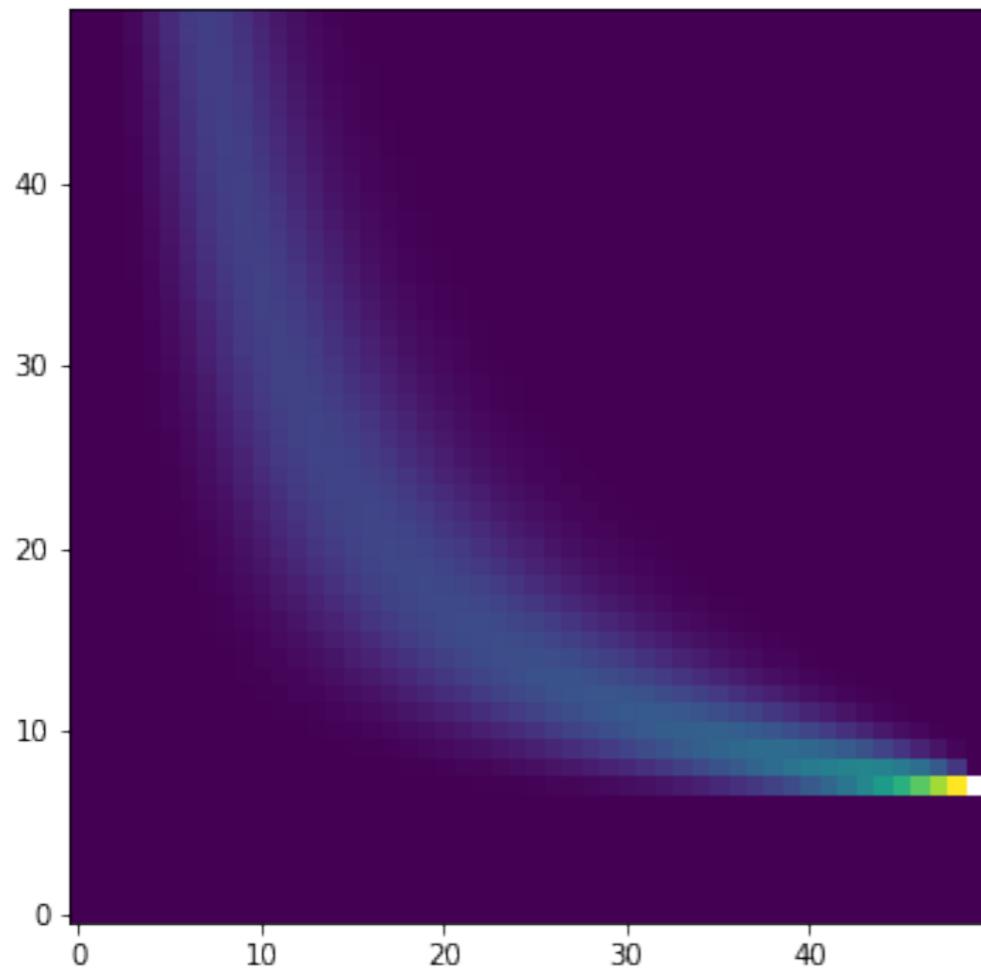
```
logp_m
```

```
<pymc3.model.LoosePointFunc at 0x125656048>
```



Coin flip in PyMC3

```
with pm.Model() as m:  
    n = pm.DiscreteUniform('n', 0, 25)  
    p = pm.Uniform('p', 0., 1., transform=None)  
    y = pm.Binomial('k', n, p, observed=7)
```



How PyMC3 use Theano

When we define a PyMC3 model, we implicitly build up a Theano function from the space of our parameters to their posterior probability density up to a constant factor.

`m.logp`

$$f: \mathbb{R} \times \mathbb{N} \rightarrow \mathbb{R}$$

`m.free_RVs`
[n, p]

<http://docs.pymc.io/theano.html>



How PyMC3 use Theano

m.logp

$$f: \mathbb{R} \times \mathbb{N} \rightarrow \mathbb{R}$$

m.free_RVs
[n, p]

n.distribution.logp

```
<bound method DiscreteUniform.logp of
<pymc3.distributions.discrete.DiscreteUniform
object at 0x125ac6358>>
```

n.logpt



How PyMC3 use Theano

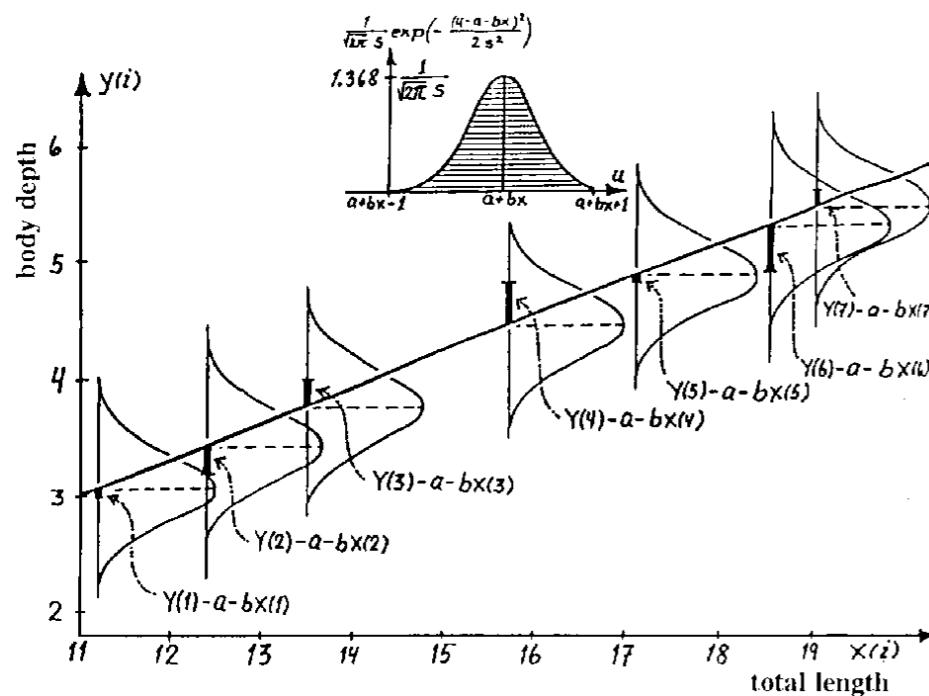
m.logp

$$f: \mathbb{R} \times \mathbb{N} \rightarrow \mathbb{R}$$

```
m.logpt??  
@property  
def logpt(self):  
    """Theano scalar of log-probability of the model"""\n    with self:  
        factors = [var.logpt for var in self.basic_RVs] +  
self.potentials  
        logp = tt.sum([tt.sum(factor) for factor in factors])  
        if self.name:  
            logp.name = '__logp_%s' % self.name  
        else:  
            logp.name = '__logp'  
    return logp
```



General linear model



ANOVA				
Source	d.f.	SS	MS	F
Treatment	$a - 1$	SS_{treat}	$\frac{SS_{\text{treat}}}{a-1}$	$\frac{MS_{\text{treat}}}{MS_{\text{error}(a)}}$
Error (a)	$N - a$	$SS_{\text{error}(a)}$	$\frac{SS_{\text{error}(a)}}{N-a}$	
Time	$t - 1$	SS_{time}	$\frac{SS_{\text{time}}}{t-1}$	$\frac{MS_{\text{time}}}{MS_{\text{error}(b)}}$
Treat x Time	$(a - 1)(t - 1)$	$SS_{\text{treat} \times \text{time}}$	$\frac{SS_{\text{treat} \times \text{time}}}{(a-1)(t-1)}$	$\frac{MS_{\text{treat} \times \text{time}}}{MS_{\text{error}(b)}}$
Error (b)	$(N - a)(t - 1)$	$SS_{\text{error}(b)}$	$\frac{SS_{\text{error}(b)}}{(N-a)(t-1)}$	
Total	$Nt - 1$	SS_{total}		

$$\mathbf{y} = \begin{pmatrix} y_{11} \\ y_{12} \\ y_{21} \\ y_{22} \\ y_{23} \\ y_{31} \\ y_{32} \end{pmatrix} = \begin{pmatrix} 1 & 1 \\ 1 & 1 \\ 1 & 2 \\ 1 & 2 \\ 1 & 2 \\ 1 & 3 \\ 1 & 3 \end{pmatrix} \begin{pmatrix} \beta_0 \\ \beta_1 \end{pmatrix} + \begin{pmatrix} \varepsilon_{11} \\ \varepsilon_{12} \\ \varepsilon_{21} \\ \varepsilon_{22} \\ \varepsilon_{23} \\ \varepsilon_{31} \\ \varepsilon_{32} \end{pmatrix} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon},$$

$$\mathbf{y} = \begin{pmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} \mu_1 \\ \mu_2 \\ \mu_3 \end{pmatrix} + \boldsymbol{\varepsilon} = \mathbf{X}_* \boldsymbol{\beta}_* + \boldsymbol{\varepsilon},$$

Generalised linear model

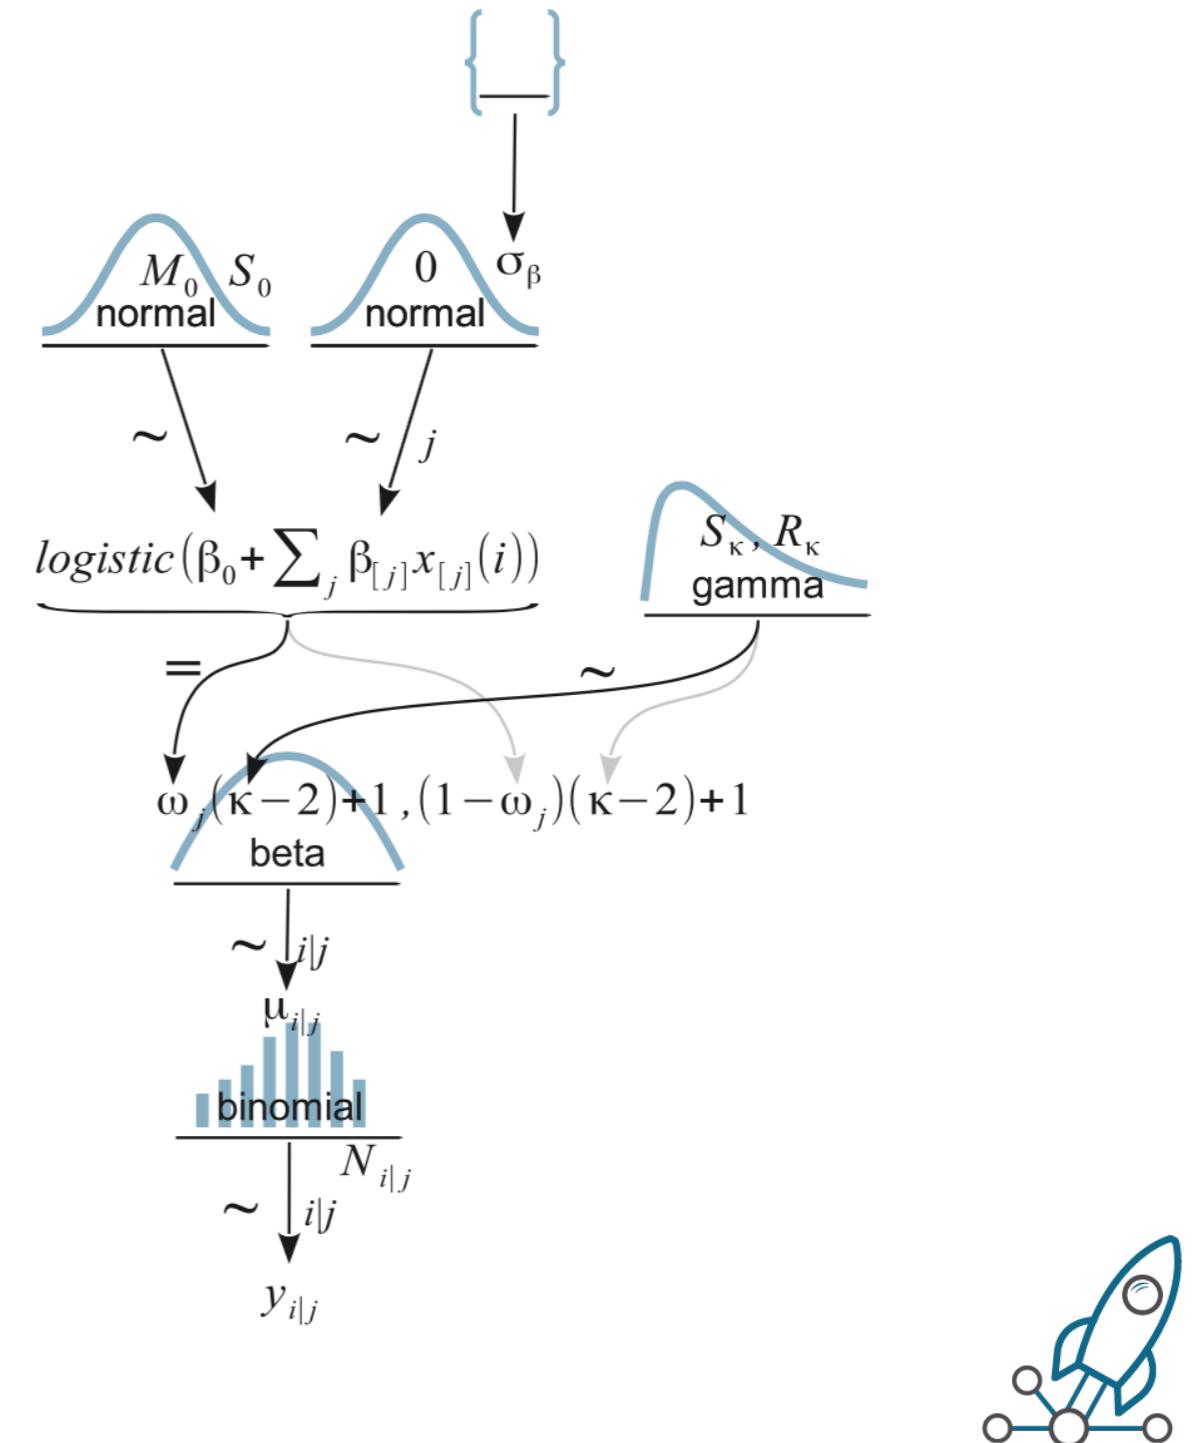
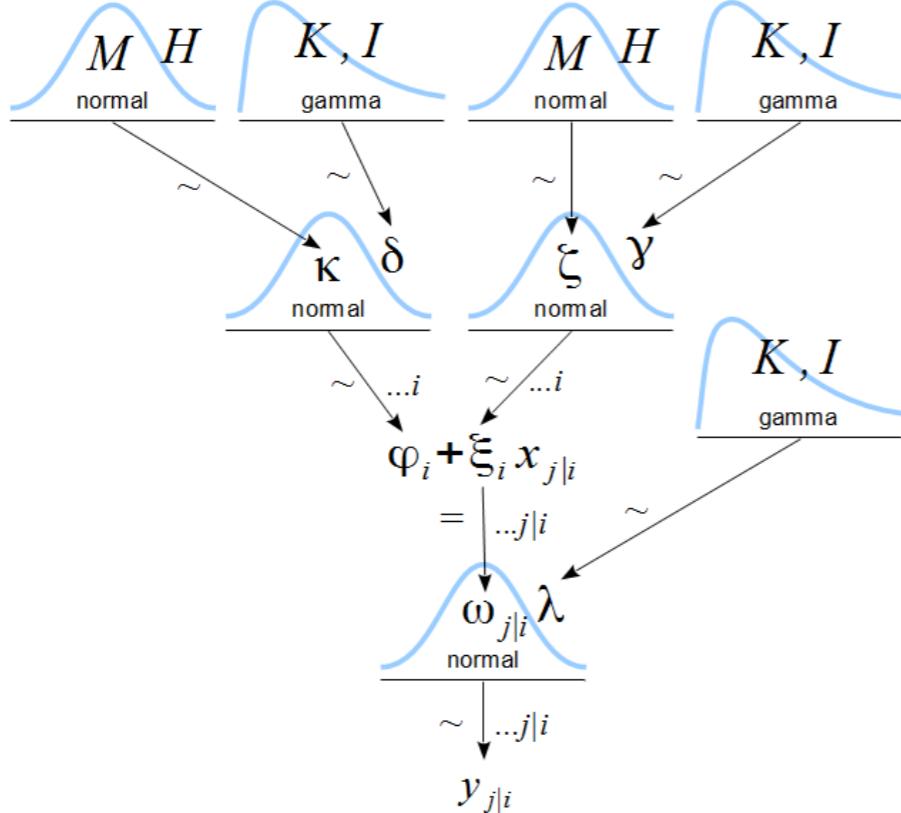
$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}, \quad E(\mathbf{y}) = \mathbf{X}\boldsymbol{\beta} = \boldsymbol{\mu}, \quad E(\boldsymbol{\varepsilon}) = \mathbf{0}, \quad \text{cov}(\mathbf{y}) = \text{cov}(\boldsymbol{\varepsilon}) = \sigma^2 \mathbf{V}.$$

outcome_i ~ Normal(μ_i, σ)

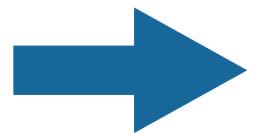
$\mu_i = \beta \times \text{predictor}_i$

$\beta \sim \text{Normal}(0, 10)$

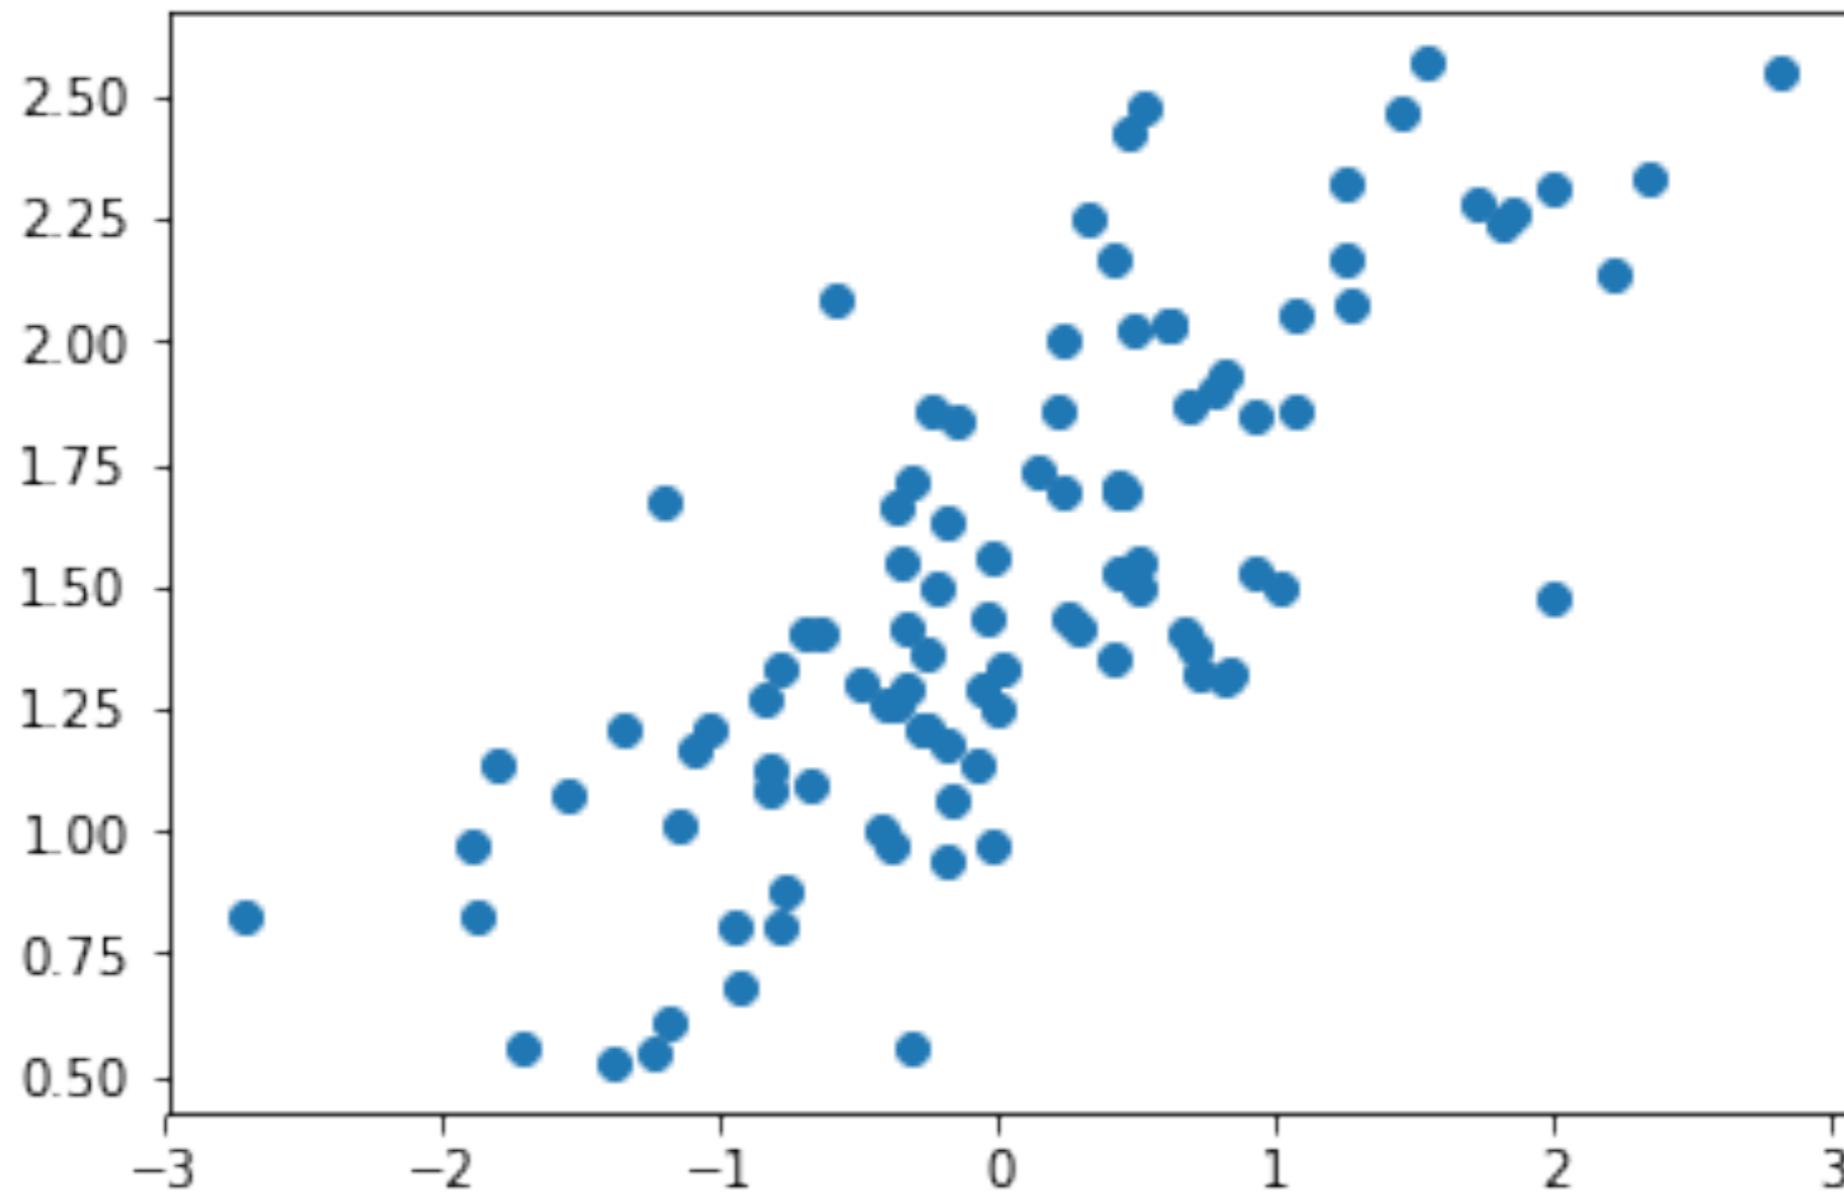
$\sigma \sim \text{HalfCauchy}(0, 1)$



Linear regression in PyMC3



Code4 - Linear_Regression.ipynb



Linear regression in PyMC3

```
with pm.Model() as m0:
```

```
    beta = pm.Normal('beta', 0., 10.)
```

```
    a = pm.Normal('a', 0., 10.)
```

```
    pm.Normal('y', x*beta+a, 1., observed=y)
```

```
with pm.Model() as m1:
```

```
    beta = pm.Flat('beta')
```

```
    a = pm.Flat('a')
```

```
    pm.Potential('logp_beta',
```

```
                 pm.Normal.dist(0., 10).logp(beta))
```

```
    pm.Potential('logp_a',
```

```
                 pm.Normal.dist(0., 10).logp(a))
```

```
    pm.Potential('logp_obs',
```

```
                 pm.Normal.dist(x*beta+a, 1.).logp(y))
```



Linear regression in PyMC3

```
with pm.Model() as m0:
```

```
...
```

```
pm.Normal('y', x*beta+a, sd, observed=y)
```

```
with pm.Model() as m1:
```

```
...
```

```
pm.Normal('eps', 0, sd, observed=y - x*beta - a)
```



Linear regression in PyMC3

```
with pm.Model() as m0:
```

```
...
```

```
    pm.Normal('y', x*beta+a, sd, observed=y)
```

```
with pm.Model() as m1:
```

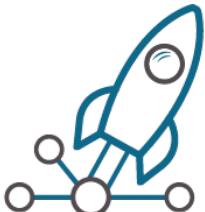
```
...
```

```
    pm.Normal('eps', 0, sd, observed=y - x*beta - a)
```

```
with pm.Model() as m1:
```

```
...
```

```
    pm.Normal('eps', 0, 1,
              observed=(y - x*beta - a)/sd)
```



Recap:

We “glue” random variables together in a conditional network to create a mapping that is the (log-) likelihood function.

The dimension of the parameter space is the same as the number of unknowns.

When we are building model, we are setting up the space (concept of coordinate system)



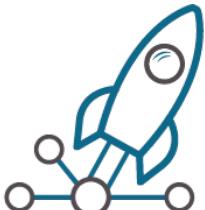
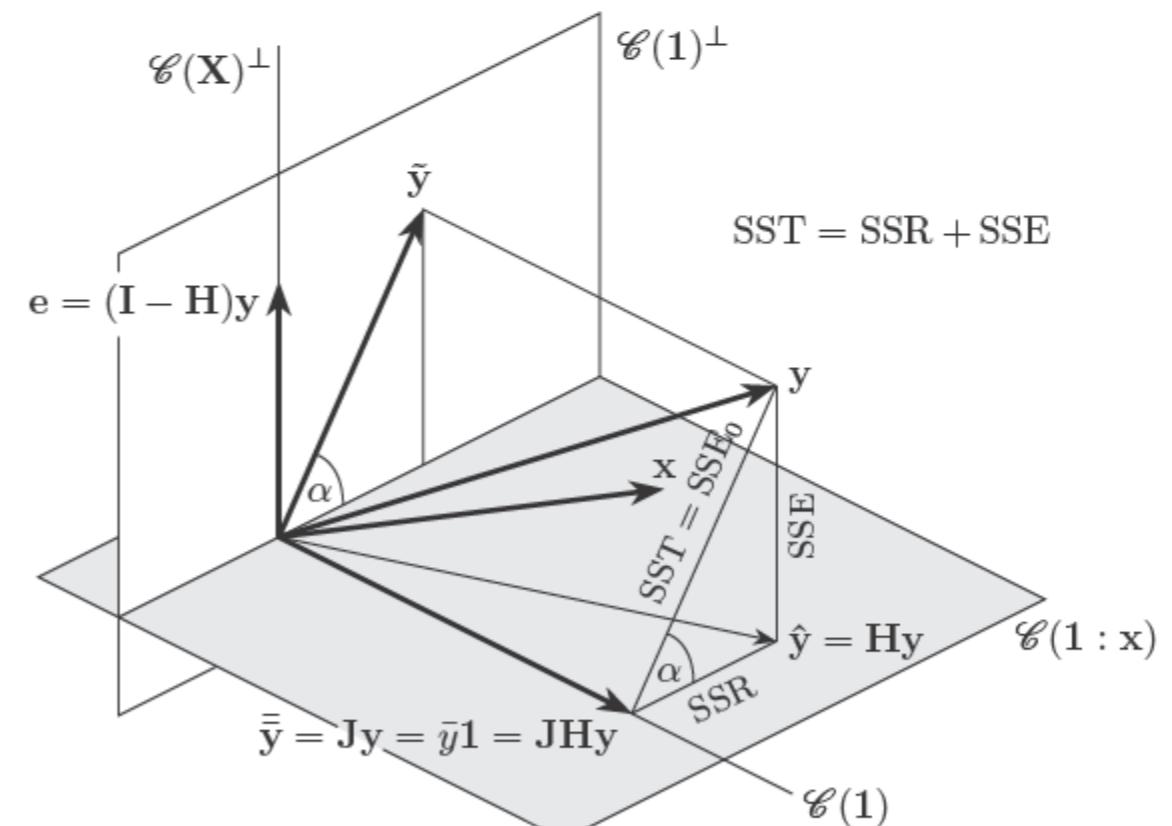
L1 and L2 regularization

Lasso Regression (Least Absolute Shrinkage and Selection Operator)

$$\sum_{i=1}^n (Y_i - \sum_{j=1}^p X_{ij}\beta_j)^2 + \lambda \sum_{j=1}^p |\beta_j|$$

Ridge regression

$$\sum_{i=1}^n (y_i - \sum_{j=1}^p x_{ij}\beta_j)^2 + \lambda \sum_{j=1}^p \beta_j^2$$



Parameterization

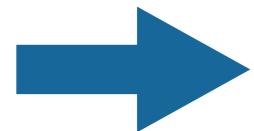
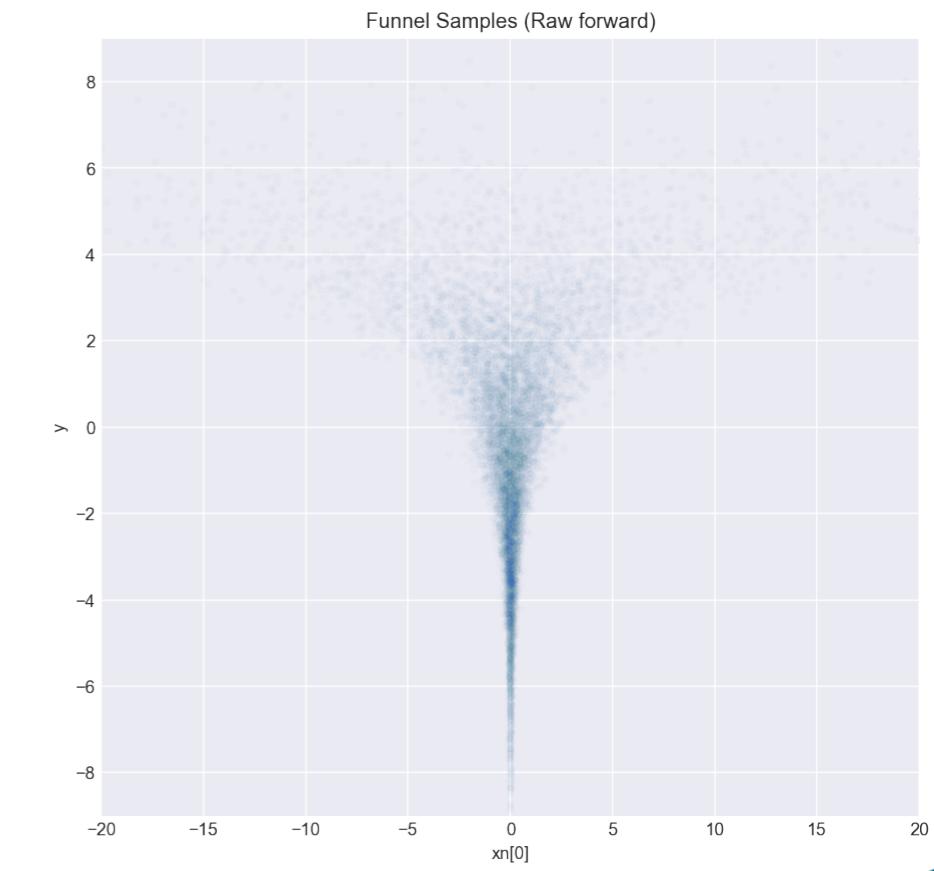
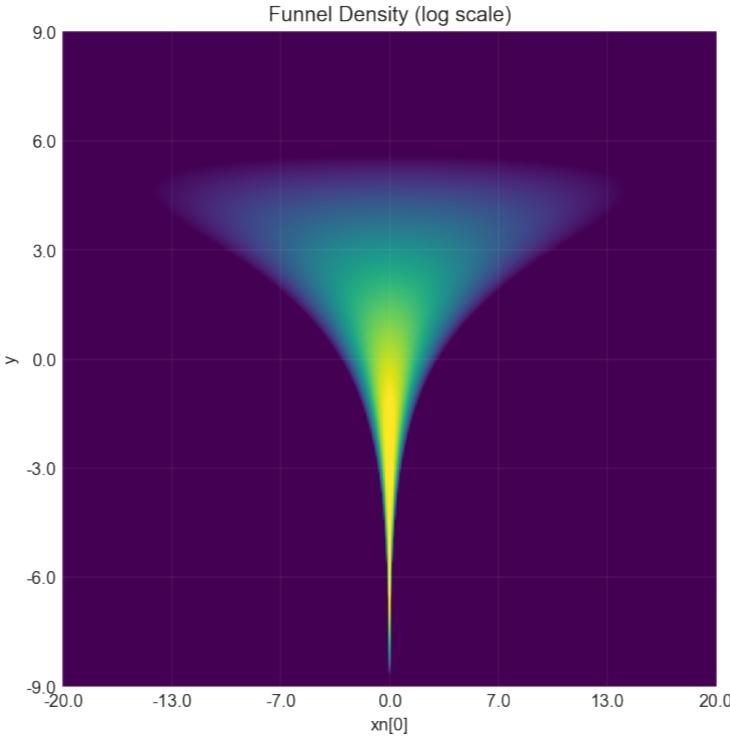
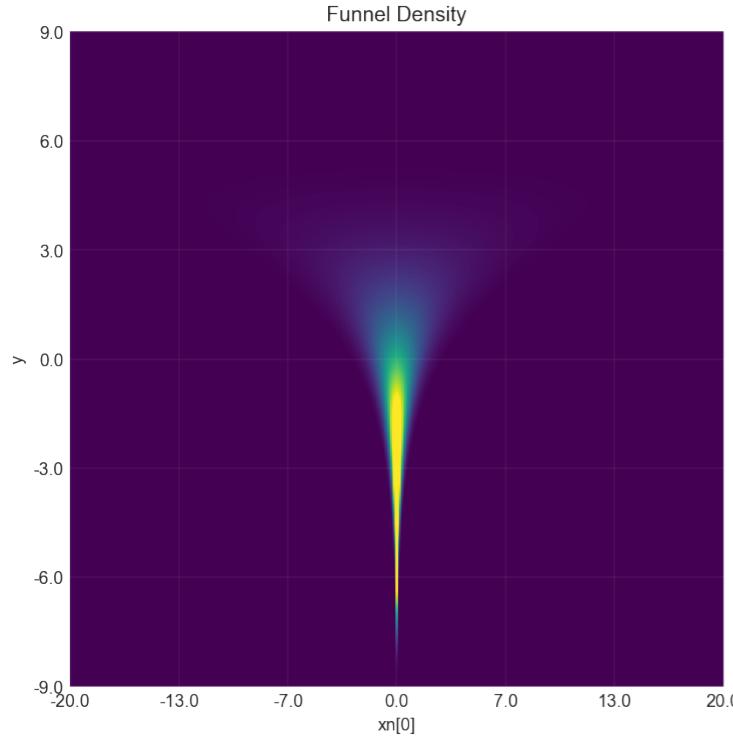
There are many ways to write down the same model:

- Make sure it is correct
- Slow model usually indicates problem
 - *The folk theorem of statistical computing: When you have computational problems, often there's a problem with your model*
- A good parameterisation should be easy to understand

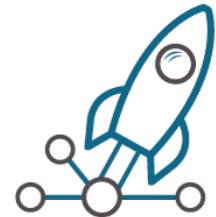


Example: Neal's Funnel

$$p(y, x_n) = \text{Normal}(y \mid 0, 3) \times \prod_{n=1}^9 \text{Normal}(x_n \mid 0, \exp(y/2))$$



Code5 - Neals_funnel.ipynb



SESSION 1: IN ALL LIKELIHOOD